

blib Code Reference v1.4

David Hobach

Contents

blib	4
Disclaimer	4
Coding Conventions	4
Library Usage	5
Dependencies	5
Imports	6
Global Variables	6
Global Aliases	7
Functions	8
args	16
Dependencies	17
Imports	17
Global Variables	17
Functions	18
arr	20
Dependencies	20
Imports	20
Functions	21
cdoc	21
Dependencies	22
Imports	22
Functions	22
daemon	26
Dependencies	27
Imports	27
Global Variables	27
Functions	27
date	30
Dependencies	30
Imports	30
Functions	30
delay	31
Dependencies	31
Imports	31

Global Variables	31
Functions	32
dmcrypt	32
Dependencies	32
Imports	33
Functions	33
flog	34
Dependencies	35
Imports	35
Global Variables	35
Functions	35
fs	39
Dependencies	39
Imports	39
Functions	39
hash	42
Dependencies	42
Imports	42
b_hash_file [file] [algorithm]	42
b_hash_str [string] [algorithm]	42
http	43
Dependencies	43
Imports	43
Global Variables	43
Functions	43
ini	44
Dependencies	45
Imports	45
Functions	45
keys	46
Dependencies	47
Imports	47
Functions	47
meta	49
Dependencies	49
Imports	49
Functions	49
multithreading/ipcm	50
Dependencies	50
Imports	50
Functions	51
multithreading/ipcv	52
Dependencies	52
Imports	52
Functions	53
multithreading/mtx	54

Dependencies	54
Imports	54
Functions	55
multithreading/multiw	57
Dependencies	58
Imports	58
Functions	58
net	59
Dependencies	59
Imports	59
Global Variables	59
Functions	59
notify	60
Dependencies	60
Imports	60
Functions	60
os/osid	61
Dependencies	61
Imports	61
Functions	61
os/qubes4/dom0	63
Dependencies	63
Imports	64
Functions	64
proc	76
Dependencies	76
Imports	76
Functions	76
sqlite3	77
Dependencies	77
Imports	77
Functions	78
str	79
Dependencies	79
Imports	79
Functions	79
tcolors	79
Dependencies	80
Imports	80
Global Variables	80
traps	80
Dependencies	80
Imports	80
Functions	80
types	81
Dependencies	81

	Imports	82
	Global Variables	82
	Functions	82
ui		84
	Dependencies	84
	Imports	84
	Functions	85
wm		85
	Dependencies	85
	Imports	85
	Functions	85
Reference List		86

blib

blib - a bash library

The basic functions which are imported by default.

Copyright (C) 2020 David Hobach LGPLv3

version: Execute `blib version` or use `b_version`.

Disclaimer

This program is free software: you can redistribute it and/or modify it under the terms of the Lesser GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the Lesser GNU General Public License for more details.

You should have received a copy of the Lesser GNU General Public License along with this program. If not, see <https://www.gnu.org/licenses/>.

The above statements apply to all modules of blib if not mentioned otherwise.

Coding Conventions

1. embrace the KISS principle
2. some general guidelines: <http://www.kfirlavi.com/blog/2012/11/14/defensive-bash-programming/>
3. use `b_[module]_[camel case function name]` to denote functions meant to be used by users of the library, `B_[module]_[upper case var name]` for global variables

4. use BLIB__[module]__[upper case var name] for global variables (to be avoided whenever possible) not meant to be used by library users (“private” variables); library users can use setters or getters
5. use the BLIB_STORE with the above naming conventions for “private” variables whenever possible
6. use blib__[module]__[camel case function name] to denote private functions not meant to be used by library users; the function name should not contain any underscores
7. the blib module itself is the only exception which can use b_, B_, blib_, BLIB_ without module name
8. prefixes such as t_, T_ and UTD_ are exclusively related to test code
9. set exit codes (!= 0 -> issue) wherever it makes sense
10. keep the global namespace clean whenever possible
11. declare -g must be used in order to allow sourcing from a function (b_import)
12. modules must check their dependencies via b_deps
13. modules should provide a header similar to that of the blib source file in order to make the documentation generation work
14. modules may be placed in subfolders of arbitrary depth
15. use 0 to indicate true and 1 to indicate false for variables; for exit codes use a non-zero exit code to indicate the number of errors
16. functions should be tagged with the following tags, if applicable:
 - @StateChanging* - the function changes the internal state of the script in a way that will not propagate to supershells (e.g. global variables) and should thus not be called from subshells (unless the user wants the state to only change in that shell)
 - @B_E* - the function uses B_E for error handling and may thus behave differently depending on the implemented error handler

Library Usage

with the default bash options:

```
source blib
b_checkVersion 1 0 || { >&2 echo "This script depends on blib (https://github.com/3hhh/blib)"
b_import [module]
```

Dependencies

```
bats
cat
dirname
find
mktemp
mv
readlink
```

rm
sort
su
whoami

Imports

no imports

Global Variables

B_TEST_MODE

blib/B_TEST_MODE

blib will set this variable to 0, if blib is running in test mode.

This variable may be used to bypass code during testing, if bats cannot test that code due to its limitations (e.g. for EXIT traps which bats uses for itself).

B_CALLER_NAME

blib/B_CALLER_NAME

Name of the executable or script as String which called blib and any child libraries.

B_ERR

blib/B_ERR

Global variable used for error handling throughout blib, cf. B_E.

It is recommended to always set an at least partially static error message on confirmed errors as variables may be empty (which would indicate “no error” for B_E).

B_RC

blib/B_RC

Can be used to set the return code for B_E in the case of an error. It defaults to 1.

This should be set to an integer value between 1 and 255. Any other value may cause undefined behaviour.

B_LIB_DIR

blib/B_LIB_DIR

Path of the blib installation directory.

B_CONF_DIR

blib/B_CONF_DIR

Path to blib configuration directory. Modules may create subfolders named by their module name there. May not exist.

B_SCRIPT

blib/B_SCRIPT

Global variable which can be used to obtain the two global variables B_SCRIPT_DIR and B_SCRIPT_NAME as follows:

```
eval "$B_SCRIPT"
```

B_SCRIPT_NAME

blib/B_SCRIPT_NAME

Path of the sourced or executed bash script executing the eval (symlinks are resolved) of B_SCRIPT.

B_SCRIPT_DIR

blib/B_SCRIPT_DIR

Name of the sourced or executed bash script executing the eval (symlinks are resolved) of B_SCRIPT.

Global Aliases

Alias expansion is automatically enabled by blib as it is required for its core functionality. So if you have strange aliases defined in your shell environment, this may cause undefined blib behaviour.

B_E

blib/B_E

The blib error handler: All blib modules use it whenever execution errors require special handling that the currently executing code cannot achieve.

Syntax:

```
B_ERR="This is an error message." ; B_E ;
```

If you need to set the return/exit code, you can do it with B_RC:

```
B_ERR="This is another error message." ; B_RC=6 ; B_E ;
```

Calling B_E means:

Check B_ERR for an error message and if there is one, handle it. It can be placed at the end of a line or on its own line. B_E will then process the error message in the way defined by the error handler (cf. b_defaultErrorHandler) and stop any further execution of code in the current context (function, script,

...) returning a non-zero exit code (1) unless the described error was fixed. In the latter unlikely case it'll let execution proceed.

More examples:

```
#handle a potentially failing command:
cmd || { B_ERR="cmd failed." ; B_E }

#capture stdout of a potentially failing command:
local out=
out="$(cmd)" || { B_ERR="cmd failed." ; B_E }

#handle multiple potentially failing commands in a try/catch like manner:
(
  set -e
  cmd1
  cmd2
  cmd3
) #NOTE: || doesn't work here!
if [ $? -ne 0 ] ; then
  B_ERR="Some commands failed."
  B_E
fi
```

The error handler can be re-defined at runtime with `b_setErrorHandler`.

Functions

`b__printStackTrace [skip level]`

blib/b__printStackTrace

[skip level]: skip that many levels of the stack trace (optional, default: 1 - skip this function call)

print the current stack trace in a human readable way

returns: stack trace with the first levels skipped as defined

`b__nop`

blib/b__nop

Do nothing.

returns: nothing; sets a zero exit code

`b__version [part]`

blib/b__version

Get the version of this blib instance.

[part]: Optional parameter defining the part of the version to retrieve (0: all as String (default), 1: major as Integer, 2: minor as Integer).

returns: blib version as string; always sets a zero exit code

b_checkVersion [minimum allowed major] [minimum allowed minor]
[maximum allowed major] [maximum allowed minor]

blib/b_checkVersion

Check whether the currently running blib instance meets the given blib version requirements.

To e.g. make sure you're on blib version 1.1 or higher, use

```
source blib
b_checkVersion 1 2 || { >&2 echo "This script depends on blib (https://github.com/3h3h3/blib)"
```

[minimum allowed major]: The lowest acceptable major version number (default: 0).

[minimum allowed minor]: The lowest acceptable minor version number (default: 0).

[maximum allowed major]: The highest acceptable major version number (default: infinite).

[maximum allowed minor]: The highest acceptable minor version number (default: infinite).

returns: Sets a zero exit code, if and only if the currently running blib instance meets the version requirements.

b_defaultErrorHandler [error out] [send err] [send stack trace]

blib/b_defaultErrorHandler

The blib default error handler.

As any error handler it must

1. handle the error message (if not the error itself) lying in B_ERR
2. not take any non-numeric arguments
3. not error out itself
4. implement the below *[error out]* as its first parameter (to make b_setBE work)
5. ideally use b_error to send error messages to the user
6. return one of the following exit codes:
 - a) 0: **if and only if** the error was fixed entirely and the caller may ignore the error (i.e. probably never)
 - b) 1: The error wasn't fixed. Functions should return to their caller indicating an error (non-zero status code). Direct shell calls will exit. B_ERR is **not** reset to blank, i.e. the next call to B_E in the same context will cause another error. The caller may use this to either *throw* the error further or handle and clear the error.
 - c) 2: Force a stop of execution in the current shell / error out.

[error out]: Whether or not to call exit after the error message handling, if the error couldn't be handled (default: 0 = always error out / call exit). If set to 1, B_E will allow e.g. functions to return to their callers.

[send err]: Whether or not to send the error message to the user (default: 0 =

send).

[send stack trace]: Whether or not to send a stack trace to the user (default: identical to [send err]).

returns: see the description above

b_setBE [error out]

blib/b_setBE

Set the [error out] behaviour of the currently configured error handler.

Contrary to `b_setErrorHandler` this function may be called by blib modules as all error handlers are required to support [error out] as parameter.

Example for switching the error out behaviour:

```
b_setBE 1
funcThatMayCallB_E #without subshell
ret=$?
b_resetErrorHandler
```

[error out]: see `b_defaultErrorHandler` (default: 0)

returns: nothing, always sets a zero exit code

@StateChanging

b_setErrorHandler [handler]

blib/b_setErrorHandler

Set the error handler for all future executions of `B_E` in the current scope.

You can do this in e.g. subshells to limit the effect.

blib modules should only use this function if absolutely necessary to temporarily modify the error behaviour whilst making sure that `b_resetErrorHandler` is called in the end. Otherwise it will prevent library users from setting the general behaviour in their scripts.

Usually you do not want to write an entirely new handler, but modify the `b_defaultErrorHandler` parameters with this setter or use `b_setBE` for that.

[handler]: Function to handle errors. See `b_defaultErrorHandler` for details.

returns: nothing

@StateChanging

b_resetErrorHandler [reset B_ERR]

blib/b_resetErrorHandler

Set the error handler to whatever it was before the last call to `b_setErrorHandler` or `b_setBE`.

[reset B_ERR]: Whether or not to also reset `B_ERR` (default: true/0).

returns: nothing, always sets a zero exit code

@StateChanging

b_getErrorHandler

blib/b_getErrorHandler

Get the currently for B_E configured error handler.

returns: the error handler function

b_silence [function] [param 1] .. [param p]

blib/b_silence

Call the given function with its parameters in the current shell context whilst suppressing all of its output to both stdout and stderr. Anything written to B_ERR however is passed to B_E (which can still write to stderr).

This function is useful when you want to keep an error message set with B_ERR, but discard everything else.

In contrast **yourfunction &> /dev/null** may also drop the error message, if you're using an error handler (see B_E) that writes to stdout or stderr.

[function]: The function to execute.

[param p]: An arbitrary number of function parameters.

returns: Sets the status code of the called function, but doesn't print anything. B_E is called on errors.

@B_E

b_defaultMessageHandler [message type] [message] [first part] [last part]

blib/b_defaultMessageHandler

Handles the given message by printing information to stdout and errors to stderr. This is the default message handler used by blib. It can be changed by b_setMessageHandler.

Message handler implementations must support at least the parameters of this function, but may add additional parameters. If an handler implementation does not support partial messages (i.e. the [first part] and [last part] parameters), it should be chained to a b_cachingMessageHandler.

[message type]: 0=informational message, 1=error message.

[message]: String representing the message.

[first part]: If set to 0, assume that the given message is the first part of an overall chain of messages (default: 0).

[last part]: If set to 0, assume that the given message is the final part of an overall chain of messages (default: 0).

returns: Nothing. Never causes errors.

b_cachingMessageHandler [message handler] [message type] [message] [first part] [last part]

blib/b_cachingMessageHandler

A message handler implementation which caches partial messages for a receiving message handler that cannot handle them until they are completed.

Use `b_initCachingMessageHandler` `&&` `b_setMessageHandler` `"b_cachingMessageHandler` [your message handler] to set your receiving message handler as message handler to cache messages for. This implementation uses the `b_getDefaultMessageHandlerIntermediate` as separator between partial messages.
[message handler]: The function implementing the receiving message handler. It must handle the [message type] and the [message].
[message type]: See `b_defaultMessageHandler`.
[message]: See `b_defaultMessageHandler`.
[first part]: See `b_defaultMessageHandler`.
[last part]: See `b_defaultMessageHandler`.
returns: Nothing. Never causes errors.

b_initCachingMessageHandler [maximum message type]

blib/b_initCachingMessageHandler

Initializes a new cache for a `b_cachingMessageHandler`. This function *must* be called at least once before using `b_cachingMessageHandler`.

[maximum message type]: Maximum message type for which to allocate caches (default: 1 -> allocate caches for 0..1/info..error).

returns: Nothing.

@B_E

b_setDefaultMessageHandlerIntermediate [intermediate]

blib/b_setDefaultMessageHandlerIntermediate

[intermediate]: String to use as intermediate between two message parts with `b_defaultMessageHandler`, i.e. the resulting message should be [part 1][intermediate][part 2]. Default: Space.

returns: Nothing.

b_getDefaultMessageHandlerIntermediate

blib/b_getDefaultMessageHandlerIntermediate

Get the currently configured intermediate string for the `b_defaultMessageHandler`.

returns: The intermediate String.

b_setDefaultMessageHandlerPrefix [message type] [prefix]

blib/b_setDefaultMessageHandlerPrefix

Set the prefix to use for the default message handler and the given message type.

[message type]: 0=informational message, 1=error message.

[prefix]: A string to prefix all messages of the given type.

returns: Nothing.

b_getDefaultMessageHandlerPrefix [message type]

blib/b_getDefaultMessageHandlerPrefix

Get the prefix used for the default message handler and the given message type.

[message type]: 0=informational message, 1=error message.

returns: The currently configured prefix.

b_info [message] [first part] [last part] [message handler param 1] ...
[message handler param n]

blib/b_info

Send an informational message to the user. The message is dispatched via the currently configured message handler (default: b_defaultMessageHandler).

[message]: to send

[first part]: If set to 0, assume that the given message is the first part of an overall chain of messages (default: 0).

[last part]: If set to 0, assume that the given message is the final part of an overall chain of messages (default: 0).

[message handler param i]: Arbitrary parameters to pass to the currently configured message handler (cf. b_setMessageHandler).

returns: Nothing, always sets a zero exit code.

b_error [message] [first part] [last part] [message handler param 1]
... [message handler param n]

blib/b_error

Send an error message to the user *without* erroring out. The message is dispatched via the currently configured message handler (default: b_defaultMessageHandler).

99.9% of all users will want to use the combination of B_ERR and B_E for proper error handling instead.

[message]: to send

[first part]: If set to 0, assume that the given message is the first part of an overall chain of messages (default: 0).

[last part]: If set to 0, assume that the given message is the final part of an overall chain of messages (default: 0).

[message handler param i]: Arbitrary parameters to pass to the currently configured message handler (cf. b_setMessageHandler).

returns: Nothing, always sets a zero exit code.

b_setMessageHandler [handler]

blib/b_setMessageHandler

Set the handler used to send messages to the user. By default, b_defaultMessageHandler is used.

[handler]: Function to handle the messages. See b_defaultMessageHandler for the requirements.

returns: Nothing.

b__getMessageHandler

blib/b__getMessageHandler

Get the currently configured message handler.

returns: The message handler.

b__enforceUser [user name]

blib/b__enforceUser

enforce that the user is the given one and if not, exit the script and set a non-zero status code

[user name]: user name to check against

returns: nothing

@B_E

b__isFunction [potential function name]

blib/b__isFunction

check whether the given function is defined

returns: zero exit code if the function is defined

b__getBlibModules

blib/b__getBlibModules

get all available blib module names as a newline-separated list

returns: all available blib module names as newline-separated list

b__listContains [list] [entry]

blib/b__listContains

check whether the given list contains the given entry

[list]: newline-separated list

[entry]: string to be found on a single line within the list (equality check)

returns: a zero exit code if the list contains the entry; a non-zero exit code otherwise

b__deps [dependency 1] ... [dependency n]

blib/b__deps

Assert that the given dependencies are met and error out with B_E otherwise. This function is meant to be used by modules or scripts to declare all of their dependencies.

[dependency i]: Command that is absolutely required to run this script.

returns: Nothing. Errors out with B_E, if dependencies are not met.

@B_E

b_import [module] [double import]

blib/b_import

Import the given module into the current context.

[module]: relative path of the module to import (relative to the blib/lib root directory)

[double import]: if set to 1, import the given module regardless of whether it was imported before (default: 0 = don't do duplicate imports)

returns: nothing, errors out if the import failed and sets a non-zero status code; if the import was successful or previously done, a zero exit code is set

@StateChanging

@B_E

b_generateStandalone [function] [module dep 1] .. [module dep n] - [function dep 1] .. [function dep d] - [function param 1] .. [function param p]

blib/b_generateStandalone

Create a standalone variant of blib in a single file running the given function when called (sourcing that file will only make the functions available) and print that file to stdout.

The current execution state is not retained.

[function]: The function to call when the generated script is executed. All script parameters when calling [output file] are passed to this function. The function must be available in the current context.

[module dep i]: Names of the modules to include in the standalone file. They do not need to be imported. They are loaded in the specified order.

[-]: Dash used as separator between the various types of arguments. If none is provided, all parameters are assumed to be modules.

[function dep j]: An arbitrary number of functions that need to be added in order to satisfy the dependencies of the function to call (e.g. if function A is meant to be called, but uses function B internally, you'll have to pass B as one of its dependencies). Dependencies that can be found in added modules should *not* be added.

[function param p]: Static parameters to add to the function as single String. Dynamic parameters should be passed to the generated script.

returns: Sets a zero exit code and prints the output file to stdout on success. May error out otherwise.

@B_E

b_execFuncInCurrentContext [function] [module dep 1] .. [module dep n] - [function param 1] .. [function param p]

blib/b_execFuncInCurrentContext

Execute the given function in the current context.

[function]: The function to execute.

[module dep i]: Names of the modules required by the function. They do not need to be imported by the function itself.

[-]: A dash as separator character between the various parameters.

[function param p]: An arbitrary number of function parameters.

returns: Whatever the executed function returns.

b_execFuncAs [user] [function] [module dep 1] .. [module dep n] -
[function dep 1] .. [function dep d] - [function param 1] .. [function
param p]

blib/b_execFuncAs

Attempt to execute the Bash function as the given user.

Whether or not this works highly depends on the underlying OS and its (sudo & su) configuration. In particular this function may cause further execution to wait for the user to type in the password of the requested user.

If the given user is identical to the current user, **b_execFuncAs** may decide to run the function in the current context. Otherwise it may run in a different process, i.e. all initialization and state may be lost. Thus make sure to create any required state within your function.

[user]: User to execute the function as (default: root).

[function]: The function to execute.

[module dep i]: Names of the modules required by the function. They do not need to be imported by the function itself. They are loaded in the specified order.

[-]: A dash as separator character between the various parameters.

[function dep j]: An arbitrary number of functions that need to be added in order to satisfy the dependencies of the function to call (e.g. if function A is meant to be called, but uses function B internally, you'll have to pass B as one of its dependencies). Dependencies that can be found in added modules should *not* be added.

[function param p]: An arbitrary number of function parameters.

returns: Whatever the executed function returns. A non-zero exit code may also indicate that the user switch didn't work. In particular **B_E** is *not* called if the executed function returns an error.

@B_E

b_isModule [module name]

blib/b_isModule

Test whether the given name represents a blib module name.

returns: sets a zero exit code if the given name is a valid module name

args

Stateful argument parser for bash.

Regular arguments and options are parsed via `b_args_parse`. Afterwards they can be retrieved via `b_args_get` and `b_args_getOption`. Options may have parameters, can be repeated and combined.

The module can also check the correctness of options. All remaining correctness checks (e.g. for regular or option parameters) are left to the user of this module.

Conventions:

- Options always start with `-` and may occur everywhere.
- Single character options may be combined, e.g. `-ajh` will be considered the same as `-a -j -h`. Long options such as `--option` cannot be combined.
Recommendation: Use long options for those options which require parameters and single letter options for everything else.
- Everything after a space-separated double dash (`--`) is not considered an option, but a regular argument.

Copyright (C) 2020 David Hobach LGPLv3
0.9

Dependencies

no dependencies

Imports

types

Global Variables

B_ARGS

args/B_ARGS

Array of regular/non-option arguments in the order of their appearance.

Instead of using it directly, it is recommended to use `b_args_get` instead. The array may be removed in future versions.

B_ARGS_OPTS

args/B_ARGS_OPTS

Map of options. `[option]_[index]` is used as key with `[index]` starting at zero. The values are the option parameters (if any). Multiple option parameters are separated by tabs.

The latter separator can be changed via `b_args_setOptionParamSeparator` before calling `b_args_parse`.

The index will only increase if options are repeated.

Instead of using it directly, it is recommended to use `b_args_getOption` instead. The array may be removed in future versions.

Functions

b_args_setOptionParamSeparator [separator]

args/b_args_setOptionParamSeparator

Set the separator for multiple option parameters.

[separator]: String to use as separator for multiple option parameters.

returns: Nothing.

b_args_getOptionParamSeparator

args/b_args_getOptionParamSeparator

Get the separator for multiple option parameters.

returns: Nothing.

b_args_init [allow flag] [option 1] [option param count 1] ... [option n] [option param count n]

args/b_args_init

Initialize the args module. It is recommended to call this function before using this module, if you want to achieve any non-default behaviour.

[allow flag]: If set to 0, assume that non-specified options do not have any parameters (default). If set to 1, enforce that only the given options are allowed and otherwise error out.

[option i]: An allowed option including its leading - prefix.

[option param count i]: Number of expected parameters for that option. The parameters must directly follow the option. If less parameters are found, the parser will error out.

returns: Nothing, always sets a zero exit code.

@StateChanging

@B_E

b_args_parse [arguments]

args/b_args_parse

Parse the given arguments.

Call **b_args_init** before parsing, if you desire any non-default parsing behaviour.

[arguments]: The arguments meant to be parsed, usually "\$@".

returns: Sets a zero exit code on success and calls **B_E** otherwise.

@StateChanging

@B_E

b_args_assertOptions [option 1] ... [option n]

args/b_args_assertOptions

Assert that the parsed options contain only the given allowed options (or less).

[option i]: Option to check against.

returns: Sets a zero exit code, if the current parsed state contains only allowed

options and errors out with B_E otherwise. Prints a list of invalid options in the error case.

@B_E

b_args_get [index] [fallback]

args/b_args_get

Get the argument at the given index.

[index]: Index of the argument to retrieve, starting at 0.

[fallback]: Value to return, if the given index was provided as empty argument or is missing. Default: empty String

returns: The argument at the given index. An *empty* (existing!) argument will cause the fallback to be returned with a zero exit code. If no argument was found at that index, a nonzero exit code is set and the fallback is returned.

b_args_getInt [index] [fallback]

args/b_args_getInt

Convenience wrapper for b_args_get that also checks the type of the argument to be an integer.

returns: See b_args_get. Also sets a nonzero exit code, if the argument is not an integer.

b_args_getCount

args/b_args_getCount

Get the number of arguments.

returns: Number of arguments.

b_args_getOption [option] [fallback] [repeat index] [parameter index]

args/b_args_getOption

Check whether the given option is set and retrieve its parameter, if it was.

[option]: String defining the option, e.g. --option or -a.

[fallback]: Value to return if the option did not have any parameter or was not set. Default: empty String

[repeat index]: Index of the option to retrieve, if the option was repeated multiple times (default: 0 = first option).

[parameter index]: Index of the option parameter to retrieve, starting at 0 (default: return all option parameters, separated by b_args_getOptionParamSeparator).

returns: The option parameter, if the option was set. The fallback is returned and a zero exit code is set, if the option was set, but an *empty* (existing!) parameter was provided. A nonzero exit code indicates that the option was not set (the fallback is still returned).

b_args_getOptionInt [option] [fallback] [repeat index] [parameter index]

args/b_args_getOptionInt

Convenience wrapper for `b_args_getOption` that also checks the type of the option to be an integer.

returns: See `b_args_getOption`. Also sets a nonzero exit code, if the option is not an integer.

b_args_getOptionCount

args/b_args_getOptionCount

Get the number of options that were set.

returns: Number of options (incl. repeated options).

b_args_getAll [exclude regex 1] ... [exclude regex n]

args/b_args_getAll

Retrieve all arguments (excl. options).

[exclude regex i]: Arguments matching any of the given regular expressions are excluded from the return value.

returns: All arguments as a single escaped string that can be used to pass them to functions or other scripts (without quotes).

b_args_getAllOptions [exclude regex 1] ... [exclude regex n]

args/b_args_getAllOptions

Retrieve all options.

[exclude regex i]: Options matching any of the given regular expressions are excluded from the return value.

returns: All options as a single escaped string that can be used to pass them to functions or other scripts (without quotes).

arr

Collection of array related functions.

Copyright (C) 2018 David Hobach LGPLv3
0.2

Dependencies

no dependencies

Imports

no imports

Functions

b_arr_join [delimiter] [array]

arr/b_arr_join

Join the given array; elements are separated with the given delimiter. The array is not checked to exist.

[delimiter]: String to use as delimiter.

[array]: Expanded array to join, e.g. “\${arr[@]}”.

returns: Joined version of the array. The exit code is always zero.

b_arr_toList [array]

arr/b_arr_toList

Create a newline-separated list from the given array.

[array]: Expanded array to join, e.g. “\${arr[@]}”.

returns: List version of the array. The exit code is always zero.

b_arr_contains [element] [array]

arr/b_arr_contains

Check whether an array contains an element.

[element]: element to check for its existence in the array

[array]: expanded array to check, e.g. “\${arr[@]}”

returns: an exit code of 0, if the element was found and 1 otherwise

b_arr_mapsAreEqual [map spec 1] [map spec 2]

arr/b_arr_mapsAreEqual

Check whether the two given maps/associative arrays are equal.

[map spec 1]: First map specification to check. Since maps cannot be passed directly to functions in Bash 4.2, you'll have to use “\$(declare -p "yourmap")” instead.

[map spec 2]: Second map specification to check.

returns: an exit code of 0, if the maps are equal and 1 otherwise; B_E is only triggered on programming errors

@B_E

cdoc

Generate code documentation in many formats (e.g. html, pdf, manpage, ...) from code comments.

Lines applicable for the documentation in your code are assumed to match static (configurable) regular expressions. These lines are then fed to pandoc in order to generate a single html page (or pdf, manpage, ...) as documentation. If no conversion is required (input format = desired output format), pandoc is bypassed.

It should be possible to use this way of generating code documentation with most programming languages (incl. bash). The defaults however are set for bash and the blib way of documenting its code, i.e. you'll have to use the getters and setters of this module if you want something different. For instance the default is to check for lines starting with `#+` (a special bash comment line) and add everything afterwards to the output documentation.

Various callback functions can be used to add content to the output of `b_cdoc_generate`. See the documentation of that function for details.

If you wish to create code documentation for your bash project in blib style, please use `../util/blib-cdoc`.

Copyright (C) 2020 David Hobach LGPLv3
0.6

Dependencies

cat
mktemp
mv
rm

Imports

fs

Functions

b_cdoc_setExtractionRegex [regex]

cdoc/b_cdoc_setExtractionRegex

Set the regular expression used to check for matching lines in code files. The first match (`${BASH_REMATCH[1]}`) is added to the documentation output.

returns: nothing

@StateChanging

b_cdoc_getExtractionRegex

cdoc/b_cdoc_getExtractionRegex

See the setter.

returns: The property that was set.

b_cdoc_setFileCallback [callback function name]

cdoc/b_cdoc_setFileCallback

Set the function to call by `b_cdoc_generate` exactly once before starting to process a source code file.

The callback function can be used to filter certain files from processing or add them as-is.

It should be declared as follows:

```
callback_function_name [file] [output format]
[file]:                The next file to process is passed here.
[output format]:      chosen output format
returns:              Nothing. Possible exit codes:
                      0 = continue normal processing (default)
                      1 = include the file as-is without any processing
                      2 = silently ignore that file / do not process it
                      other = abort all further processing with an error
```

returns: nothing

@StateChanging

b_cdoc__getFileCallback

cdoc/b_cdoc__getFileCallback

See the setter.

returns: The property that was set.

b_cdoc__setDocumentBeginCallback [callback function name]

cdoc/b_cdoc__setDocumentBeginCallback

Set the function to call by `b_cdoc__generate` exactly once right before it starts generating the output document.

The callback function should be declared as follows:

```
callback_function_name [document output file] [document output format]
[document output file]: path to the document output file
                       (may not exist and should not be written to)
[document output format]: chosen output format
returns:               whatever should be added at the beginning of the output document;
                       a non-zero exit code will abort further processing
```

returns: nothing

@StateChanging

b_cdoc__getDocumentBeginCallback

cdoc/b_cdoc__getDocumentBeginCallback

See the setter.

returns: The property that was set.

b_cdoc__setPostProcessingCallback [callback function name]

cdoc/b_cdoc_setPostProcessingCallback

Set the function to call by `b_cdoc_generate` each time a code file was fully processed.

The callback function should be declared as follows:

```
callback_function_name [processed input] [input file] [document output format]
[processed input]: Everything that was found to match the
                    extraction regex in the [input file] by b_cdoc_generate.
[input file]:       The original input file.
[document output format]: chosen output format
returns:           whatever should be added to the output document for the
                    given input file (usually the processed input or some filtered
                    version of it); a non-zero exit code will abort further processing
```

returns: nothing

@StateChanging

b_cdoc_getPostProcessingCallback

cdoc/b_cdoc_getPostProcessingCallback

See the setter.

returns: The property that was set.

b_cdoc_setDocumentEndCallback [callback function name]

cdoc/b_cdoc_setDocumentEndCallback

Set the function to call by `b_cdoc_generate` exactly once right after it generated the output document.

The callback function should be declared as follows:

```
callback_function_name [document output file] [document output format]
[document output file]: path to the document output file
                        (may not exist and should not be written to)
[document output format]: chosen output format
returns:               whatever should be added to the end of the output
                        document; a non-zero exit code will abort further processing
```

returns: nothing

@StateChanging

b_cdoc_getDocumentEndCallback

cdoc/b_cdoc_getDocumentEndCallback

See the setter.

returns: The property that was set.

b_cdoc__setBlockCallback [callback function name]

cdoc/b_cdoc__setBlockCallback

Set the function to call by `b_cdoc__generate` each time it hits a block of matching comments.

The callback function should be declared as follows:

```
callback_function_name [block] [input file] [document output format]
[block]:               The full block of documentation that was identified.
[block counter]:       Number of blocks previously seen.
[input file]:          The original input file.
[document output format]: chosen output format
returns:               whatever should be added instead of the given block;
                       a non-zero exit code will abort further processing
```

returns: nothing

@StateChanging

b_cdoc__getBlockCallback

cdoc/b_cdoc__getBlockCallback

See the setter.

returns: The property that was set.

b_cdoc__generate [input files] [output file] [output format] [additional pandoc options]

cdoc/b_cdoc__generate

Generate a documentation file from the given list of input files or directories.

The concept is really simple: Each *block* of documentation will trigger the `b_cdoc__getBlockCallback` exactly once and you may add additional parsing logic on a block-wise level.

The following example will call the block callback twice (once with three `block 1` lines and once with two `block 2` lines):

```
#+block 1
#+block 1
#+block 1

#say hello
echo "hello world"
#+block 2
#+block 2
```

The other callback functions may be used for further processing.

[input files]: Newline-separated list of files or directories to generate the documentation from. The given order is respected; directories are recursively searched

for files. It is currently assumed that these files are encoded in UTF-8.
[output file]: Path to the documentation file to generate. Should not exist.
[output format]: The target format of the documentation to generate. See pandoc for a list of available output formats. If none is specified, pandoc is bypassed and the input format is chosen as output format. Passing “pandoc” will let pandoc decide based on the extension of the output file.
[additional pandoc options]: All remaining parameters will be directly passed to pandoc. If none are provided, -s is implicitly added as default.
returns: Sets a non-zero exit code and exits the script on errors. Output from pandoc and other calls may be printed. Otherwise nothing is returned.
@B_E

b_cdoc_generateBlibStyle [input files] [output file] [output format]
[delete existing]

cdoc/b_cdoc_generateBlibStyle

A convenience wrapper for b_cdoc_generate which sets various reasonable parameters depending on the output format.

[input files]: see b_cdoc_generate

[output file]: where to write the generated output documentation to

[output format]: currently one of raw|html|pdf|man is supported (default: raw)

[delete existing]: whether or not to delete previously created output files (default: true/0); if set to false (1), the function will error out if a previously created file was found

returns: full path to the created documentation file on success; otherwise the function may error out

@B_E

Callback Functions

b_cdoc_cbPrintNewline

cdoc/b_cdoc_cbPrintNewline

Prints a newline character.

returns: nothing

b_cdoc_cbPrintFirstParam [param]

cdoc/b_cdoc_cbPrintFirstParam

Prints the first parameter.

[param]: The parameter to print.

returns: the first parameter

daemon

Module providing access to a single background process providing some service (daemon). Exiting the foreground control process will *not* terminate the

background process. Attempting to start multiple background daemons will be prevented in a thread-safe way.

Each background process is assumed to implement a `daemon_main` function and must be identified by a unique String.

If you need to exchange data with the background service, please have a look at the `multithreading/ipcv` or `multithreading/ipcm` modules.

Known Issue:

Since the daemon cannot detach from its session via `setsid` in `bash`, the daemon will remain in the process group of its parent. So killing the parent with e.g. `Ctrl-C` will cause the parent group including the daemon to terminate. As a workaround, users may either ignore `SIGINT` requests or make sure the parent exits as soon as possible (recommended).

Copyright (C) 2019 David Hobach LGPLv3

0.4

Dependencies

kill
mkdir
mktemp
rm
umask

Imports

`multithreading/mtx`
`proc`

Global Variables

`B_DAEMON_ID`

daemon/B_DAEMON_ID

Contains the ID of the daemon, if and only if the current process is the daemon process.

Functions

`b_daemon_init` [`quiet flag`] [`main name`] [`stdout file`] [`stderr file`]
[`umask setting`]

daemon/b_daemon_init

Init the module paramters. It is necessary to call this method *before* using any other of this module unless you want to use the default paramters.

[`quiet flag`]: If set to 0 (default), don't print anything to `stdout` during the execution of `start|stop|restart|status`. Otherwise use `b_info` to print informational

messages.

[main name]: Name of the main loop function to execute in the background process. Returning from that function will exit the background process. Default: **daemon_main**

[stdout file]: Where the background process should write its stdout stream to (default: /dev/null). Lines are appended.

[stderr file]: Where the background process should write its stderr stream to (default: /dev/null). Lines are appended.

[umask settings]: The umask settings to apply to the daemon (default: 0).

returns: Nothing.

@StateChanging

b_daemon_start [id] [arg 1] ... [arg n]

daemon/b_daemon_start

Start the background process.

If you need to start it as a different user, simply run this function as a different user with e.g. (b_execFuncAs)[#b_execFuncAs]. Please keep in mind that control processes must have the permission to send signals to the daemon PID though.

[id]: Unique identifying String of the daemon to distinguish it from others.

[arg i]: An arbitrary number of arguments which can be passed to the main loop.

returns: Sets a zero exit code on success. Otherwise sets a non-zero exit code. In particular B_E is called, if the daemon is already running.

@B_E

b_daemon_stop [id] [termination signal] [kill timeout]

daemon/b_daemon_stop

Stop the background process.

[id]: Unique identifying String of the daemon to distinguish it from others.

[termination signal]: A number or string specifying the signal to send to the daemon (default: 15 / SIGTERM). See `kill -l` for an overview.

[kill timeout]: Time in seconds after which the background process will be killed, if it remains unresponsive to the termination signal (default: 0 = wait indefinitely).

returns: Sets a zero exit code, if the daemon terminated by itself. An exit code of 2 indicates that the daemon had to be killed. An exit code of 3 means that it wasn't running. B_E is called on unexpected errors.

@B_E

b_daemon_restart [id] [termination signal] [kill timeout] [arg 1] ... [arg n]

daemon/b_daemon_restart

Restart the background process.

[id]: Unique identifying String of the daemon to distinguish it from others.

[termination signal]: See `b_daemon_stop`.
[kill timeout]: See `b_daemon_stop`.
[arg i]: An arbitrary number of arguments which can be passed to the main loop.
returns: See `b_daemon_start`.
@B_E

b_daemon_statusPid [id]

daemon/b_daemon_statusPid

Check the status of the background process.

Doesn't print informational messages to stdout.

[id]: Unique identifying String of the daemon to distinguish it from others.

returns: The PID and sets a zero exit code, if the daemon is running and a non-zero exit code otherwise. `B_E` is only called on exceptional errors.

@B_E

b_daemon_status [id]

daemon/b_daemon_status

Check the status of the background process and print informational messages to stdout (if configured).

[id]: Unique identifying String of the daemon to distinguish it from others.

returns: Sets a zero exit code, if the daemon is running and a non-zero exit code otherwise. `B_E` is only called on exceptional errors.

@B_E

b_daemon_getPid [id]

daemon/b_daemon_getPid

Get the process ID of the background process.

[id]: Unique identifying String of the daemon to distinguish it from others.

returns: The process ID and sets a zero exit code, if it could be obtained. Please note that the process may be dead anyway (use `b_daemon_statusPid` for that). Otherwise a non-zero exit code is set. `B_E` is only called on exceptional errors.

@B_E

b_daemon_sendSignal [id] [signal]

daemon/b_daemon_sendSignal

Send a signal to the daemon.

For termination signals, please use `b_daemon_stop` instead.

[id]: Unique identifying String of the daemon to distinguish it from others.

[signal]: A number or string specifying the signal to send to the daemon. See `kill -l` for an overview.

returns: Sets a zero exit code, if the daemon was running and a non-zero exit

code otherwise. B_E is only called on exceptional errors.

@B_E

date

Collection of date and time related functions.

Copyright (C) 2018 David Hobach LGPLv3

0.3

Dependencies

date

Imports

no imports

Functions

b_date_add *date* [**time**] [**unit**] [**format**] [**utc flag**]

date/b_date_add

Add the given number of seconds to the given date.

[date]: date to add seconds to; the format must be understood by the Unix date utility

[time]: amount of time to add

[unit]: unit of the time to add, may be one of d (days), h (hours) m (minutes), s (seconds, default)

[format]: output format of the date, in Unix date notation (default: use the localized output)

[utc flag]: if set to 0, use UTC as time zone if not specified for the input *and* use it for the output (default: 1 = local time zone)

returns: The input date with the given number of seconds added, in the requested format; returns a non-zero exit code on errors.

@B_E

b_date_addDays *date* [**days**] [**format**] [**utc flag**]

date/b_date_addDays

Convenience wrapper to b_date_add with days.

[date]: See b_date_add.

[days]: Number of days to add.

[format]: See b_date_add.

[utc flag]: See b_date_add.

returns: See b_date_add.

@B_E

b_date_diff [date 1] [date 2] [unit]

date/b_date_diff

Get the amount of time between the two dates, i.e. [date 2] - [date 1].

[date 2], [date 1]: the two dates to subtract; the time part is assumed to be identical if not specified within the dates

[unit]: unit of the result, may be one of d (days), h (hours) m (minutes), s (seconds, default)

returns: The amount of time between the given two dates [date 2] - [date 1], rounded down. Returns a non-zero exit code on errors.

@B_E

b_date_getFileModAge [file] [unit]

date/b_date_getFileModAge

Get the time that passed since the last modification of the file.

[file]: Full path to the file to check.

[unit]: unit of the time to retrieve, may be one of d (days), h (hours) m (minutes), s (seconds, default)

returns: The amount of time in the given unit since the last modification. May be rounded down. Sets a non-zero exit code on errors.

@B_E

delay

Simplistic module to delay commands to a future time.

Requires polling.

All timestamps in this module must be integers and a larger timestamp must denote a time *after* a smaller timestamp.

Copyright (C) 2019 David Hobach LGPLv3

0.3

Dependencies

no dependencies

Imports

no imports

Global Variables

B_DELAY_EXECUTED

delay/B_DELAY_EXECUTED

The number of commands executed during the last invocation of b_delay_execute.

Functions

b_delay_to [timestamp] [command]

delay/b_delay_to

Delay the given command to be executed at the given time.

[timestamp]: An integer timestamp (e.g. \$SECONDS, Unix timestamp in s/ms/ns, ...).

[command]: The command to execute at the given time.

returns: Nothing.

@StateChanging

b_delay_execute [timestamp]

delay/b_delay_execute

Execute all commands which are due at the given time.

[timestamp]: Integer timestamp representing the current point in time.

returns: Nothing. The exit code is equal to the number of commands with a non-zero exit code. B_DELAY_EXECUTED is updated with the number of commands executed.

@StateChanging

b_delay_getCommandAt [timestamp]

delay/b_delay_getCommandAt

Get the set of commands to be executed at the given point in time.

[timestamp]: Integer timestamp denoting the time for which to retrieve the commands.

returns: The commands to execute at that time.

dmccrypt

Abstraction layer for cryptsetup / dm-crypt.

Features:

- automatic management of dm-crypt devices
- password support for non-tty environments

Copyright (C) 2020 David Hobach LGPLv3

0.6

Dependencies

dirname

head

mkdir

readlink

Imports

hash
ui

Functions

b_dmccrypt_init [ui mode]

dmccrypt/b_dmccrypt_init

Initialize this module. This function *must* be called at least once before using any of the other functions.

[ui mode]: How to request a password from the user: auto|gui|tty (default: auto).

returns: Nothing.

@StateChanging

@B_E

b_dmccrypt_getMapperName [path]

dmccrypt/b_dmccrypt_getMapperName

Get the name of the dm-crypt mapper for a given path.

[path]: Full path to an encrypted file.

returns: A mapper name. This doesn't necessarily mean that the encrypted container is open. Use `b_dmccrypt_isOpen` for that.

@B_E

b_dmccrypt_createLuks [path] [size] [fs type] [entropy source] [password prompt] [dm-crypt option 1] ... [dm-crypt option n]

dmccrypt/b_dmccrypt_createLuks

Create an encrypted luks container file at the given location.

This function may request a password from the user and usually requires root access rights.

[path]: Full path to the encrypted file to create.

[size]: Filesystem size in bytes to create. Supported suffixes: b 512, kB 1000, K 1024, MB 10001000, M 10241024, GB 100010001000, G 102410241024

[fs type]: Filesystem to create inside the encrypted container. If none is specified (default), no file system is created.

[entropy source]: Source of entropy to use for the file setup (default: /dev/urandom).

[password prompt]: Prompt string to ask the user for his password (optional).

[dm-crypt option i]: These options are directly passed to `cryptsetup`.

returns: Sets a zero exit code on success and errors out with B_E otherwise.

@B_E

b_dmccrypt_open [path] [mount point] [output var] [password prompt] [dm-crypt option 1] ... [dm-crypt option n]

dmccrypt/b_dmccrypt_open

Open/Decrypt the given container and optionally mount it.

[path]: Full path to the encrypted file.

[mount point]: Where to mount the decrypted data (optional). If no mount point is specified, it will not be mounted.

[output var]: The name of the variable to write the created device to (optional).

[password prompt]: Prompt string to ask the user for his password (optional).

[dm-crypt option i]: These options are directly passed to **cryptsetup**.

returns: Sets a zero exit code on success.

@B_E

b_dmccrypt_close [path] [dm-crypt option 1] ... [dm-crypt option n]

dmccrypt/b_dmccrypt_close

Close the given encrypted container.

[path]: Full path to the encrypted file.

[dm-crypt option i]: These options are directly passed to **cryptsetup**.

returns: A zero exit code on success. The exit code may also be zero for non-existing or already closed containers.

@B_E

b_dmccrypt_isOpen [path]

dmccrypt/b_dmccrypt_isOpen

Check whether the given encrypted container is open (not necessarily mounted).

[path]: Full path to the encrypted file.

returns: Sets a zero exit code, if and only if the container is open. B_E is only called for exceptional errors.

@B_E

flog

Flexible log writer for bash.

Features:

- arbitrary output support (files, network streams, stdout, stderr, ...) in a user-defined format
- optional log file reduction
- optional thread safety
- support for partial messages

In order to log to the system log, please use the **logger** command instead. This library is mostly meant for application logs handled in a more custom manner.

Exact format of log entries:

[header] [message]

```
[header]: Can be arbitrarily defined in the
          respective callback function. If nothing
          is defined, the below default header
          is used:
[default header] = '[default date] '
[default date]: current date in the format as used
                by date +"%F %T %Z" (the format can be changed)
```

Copyright (C) 2020 David Hobach LGPLv3
0.6

Dependencies

cat
date
mkdir
mktemp
readlink
rm
tail

Imports

fs
hash
multithreading/mtx

Global Variables

B_FLOG_SEV

flog/B_FLOG_SEV

Global map for human readable severities which may be used by users of this script.

It was inspired by the severities of RFC5424.

Currently supported values: emergency|alert|critical|crit|error|err|warning|warn|notice|informational|info|debug

Functions

b_flog_printSeverity [severity]

flog/b_flog_printSeverity

[severity]: see b_flog_init

Print the given severity in a way for logging. This function is meant to be used as building block for header functions.

returns: a printed version of the given severity for logging

b_flog_close

flog/b_flog_close

close the currently open log; is automatically called, but users may want to call it themselves to force the respective file descriptor to be closed before the program is ended

returns: nothing

@StateChanging

b_flog_init [log file name] [header callback function] [log reduction lines] [thread safe] [intermediate]

flog/b_flog_init

Initialize this log writer. This function **must** be called before any others.

[log file name]: name of the log file to write to; special files such as /dev/stdout, /dev/stderr (default), /dev/tcp, /dev/udp are supported if your bash version supports them; the file doesn't need to exist

[header callback function]: optional name of the function to be called whenever a new log entry is generated; the function must be defined as follows:

[header callback function] [severity]

[severity]: see [b_flog_log](#b_flog_log)

returns: the full header meant to be used for the current moment in time with the given severity (without knowing the message details) and sets a non-zero exit code on errors; errors may cause the message to be logged without header

[log reduction lines]: if set to a positive integer, reduce the log file approximately to that number of lines during logging (default: 3000) - see b_flog_setLogReductionLinesApprox for details; this option has no effect on non-file outputs (stdout, network output, ...)

[thread safe]: Whether calls to b_flog_log should be thread safe (0) or not (default: 1 = not thread safe).

[intermediate]: String to use as intermediate separator when chaining partial log messages (default: b_getDefaultMessageHandlerIntermediate).

returns: sets a non-zero exit code on errors and may exit the script

@StateChanging

@B_E

b_flog_log [message] [severity] [first part] [last part]

flog/b_flog_log

Log the given message with the given optional severity.

If the [thread safe] variant was chosen, may wait for other log sources to write their message first.

[message]: message to log

[severity]: users may pass arbitrary numbers or even Strings here, but it is recommended to stick to the priorities defined in \$BLIB_FLOG_SEV (default:

`#{B_FLOG_SEV["info"]})`

[first part]: If set to 0, assume that the given message is the first part of an overall chain of messages (default: 0).

[last part]: If set to 0, assume that the given message is the final part of an overall chain of messages (default: 0). Please note that other threads may be blocked from writing to the output, if the last message of a chain was not yet received.

returns: Sets a non-zero exit code on errors. `B_E` is only called, if logging failed entirely.

@B_E

b_flog_messageHandler [message type] [message] [first part] [last part]

flog/b_flog_messageHandler

A message handler implementation that handles messages by logging them via `b_flog_log`.

Issues with the logging system itself (e.g. log file not writable) are written to `stderr`.

If you don't want to log all messages and/or handle some of them differently, you can simply write a wrapper for this function.

[message type]: See `b_defaultMessageHandler`.

[message]: See `b_defaultMessageHandler`.

[first part]: See `b_defaultMessageHandler`.

[last part]: See `b_defaultMessageHandler`.

returns: Nothing. Never causes errors.

b_flog_getDateFormat

flog/b_flog_getDateFormat

Get the date format used for the header by this log writer (see "man date" for explanations).

returns: see above

b_flog_setDateFormat [format string]

flog/b_flog_setDateFormat

Set the date format used for the header by this log writer (see "man date" for explanations).

returns: nothing

@StateChanging

b_flog_getLogReductionLinesLowerBound

flog/b_flog_getLogReductionLinesLowerBound

Get the number of lines that the log file will at least have after a log file reduction.

returns: see above

b_flog_getLogReductionLinesUpperBound

flog/b_flog_getLogReductionLinesUpperBound

Get the maximum number of lines that the log file will have before it is reduced.

returns: see above

b_flog_setLogReductionLinesLowerBound [bound]

flog/b_flog_setLogReductionLinesLowerBound

Set the number of lines that the log file will at least have after a log file reduction.

[bound]: number of lines to use for that bound

returns: nothing

@StateChanging

b_flog_setLogReductionLinesUpperBound

flog/b_flog_setLogReductionLinesUpperBound

Set the maximum number of lines that the log file will have before it is reduced.

[bound]: number of lines to use for that bound

returns: nothing

@StateChanging

b_flog_setLogReductionLinesApprox [line count]

flog/b_flog_setLogReductionLinesApprox

Set the number of average number of lines that the log file should have; counts
<= 0 indicate no limit.

[line count]: reduce the log after reaching 1.2*[line count] lines to 0.8*[line count]
lines

returns: nothing

@StateChanging

b_flog_getHeaderFunction

flog/b_flog_getHeaderFunction

Get the name of the header callback function that is used.

returns: see above

b_flog_setHeaderFunction [header function]

flog/b_flog_setHeaderFunction

Set the name of the header callback function to be used.

[header function]: name of the header function to use

returns: nothing

@StateChanging

Header Functions

`b_flog_defaultHeader [severity]`

flog/b_flog_defaultHeader

Default header callback function used with `b_flog_init`.

[severity]: the default header ignores the severity

returns: the default header meant to be used for the current moment in time

`b_flog_headerDateSeverity [severity]`

flog/b_flog_headerDateSeverity

An alternative to the default header callback function which appends the severity to the default header.

[severity]: see `b_flog_init`

returns: the default header with the severity appended

`b_flog_headerDateScriptSeverity [severity]`

flog/b_flog_headerDateScriptSeverity

An alternative to the default header callback function which appends the calling script and the severity to the default header.

[severity]: see `b_flog_init`

returns: the default header with the calling script and severity appended

fs

Collection of file and file system related functions.

Copyright (C) 2018 David Hobach LGPLv3

0.3

Dependencies

findmnt

head

mktemp

mount

rm

stat

sync

wc

Imports

no imports

Functions

`b_fs_isEmptyDir [dir]`

fs/b_fs_isEmptyDir

Check whether the given directory is empty or non-existing. It is not checked whether the passed parameter is a file preventing a directory from being created.

[dir]: full path to the directory to check

returns: a zero exit code if the directory does not exist or is empty

b_fs_getLineCount [file]

fs/b_fs_getLineCount

Get the number of lines of the given file.

[file]: full path to a file

returns: the number of lines; a non-zero exit code is set on errors

@B_E

b_fs_waitForFile [file] [maximum time]

fs/b_fs_waitForFile

Sleep until the given file appears. The check interval is 1s.

[file]: full path to the file or directory to wait for

[maximum time]: maximum time in s to wait for the file to appear (default: forever)

returns: Sets a zero exit code if the file appeared and a non-zero exit code on a timeout.

b_fs_getMountpoints [device]

fs/b_fs_getMountpoints

Get all mountpoints for the given device.

[device]: Full path to the device (incl. /dev/) for which to obtain the mountpoints.

returns: A newline-separated list of mountpoints where the given device is mounted to. Sets a non-zero exit code if no such mountpoints were found.

b_fs_mountIfNecessary [device] [mount point] [enforce] [mount options]

fs/b_fs_mountIfNecessary

Mount the given device if it isn't already mounted.

[device]: Full path to the device (incl. /dev/) to mount.

[mount point]: Full path where to mount the device. If no mount point is specified, a /tmp/ mount point is chosen. Non-existing directories are created. Is ignored if another mount point already exists.

[enforce]: If set to 0, enforce the given mount point to be used in addition to potentially existing ones (default: 1). Mount options are ignored.

[mount options]: Options to pass to **mount**, if it needs to be executed (default: none).

returns: The chosen mount point or a newline-separated list of existing mount

points on success; sets a non-zero exit code on failure.

@B_E

b_fs_createLoopDeviceIfNecessary [file]

fs/b_fs_createLoopDeviceIfNecessary

Create a loop device for the given file if no old one exists. Usually requires root access rights.

[file]: File for which to create a loop device.

returns: Created loop device or previously used one (incl. /dev/). Sets a non-zero exit code, if no device could be created.

@B_E

b_fs_removeUnusedLoopDevice [device|file]

fs/b_fs_removeUnusedLoopDevice

Remove a loop device, if and only if it is unused by the operating system. Otherwise mark it for removal once it becomes unused. Usually requires root access rights.

[device|file]: Full path (incl. /dev/) to the loop device or to the backing file.

returns: Sets a zero exit code, if the loop device was not used and could thus be successfully removed *or* does not exist. A nonzero exit code indicates that the device is still being used. B_E is called on other errors.

@B_E

b_fs_parseSize [string] [check flag]

fs/b_fs_parseSize

Parse human-readable file system sizes that include units.

[string]: A string denoting a file system size of the format [number][unit]. Unit may be one of KB 1000, K 1024, MB 10001000, M 10241024, GB 100010001000, G 102410241024, and so on for T, P. If no unit is provided, the number is assumed to denote bytes. The number must be an integer.

[check flag]: Check whether the result makes sense (default: 0/true). This check will make integer overflows less likely.

returns: The respective number of bytes meant with the given string. B_E is called on parsing errors.

@B_E

b_fs_removeWithOverwrite [file] [randomness source]

fs/b_fs_removeWithOverwrite

Overwrite the given file with random data, then remove it. This is meant to prevent a potential reconstruction of the file after its removal.

Warning: The reconstruction may still work on some types of file systems or physical storage systems (e.g. on flash disks).

[file]: Full path to the file to remove. Directories are currently not supported.

[randomness source]: Device to use as source of random data (default: /dev/urandom).

returns: Nothing. B_E is called on errors.

@B_E

b_fs_enumerate [path list] [nonexisting]

fs/b_fs_enumerate

Enumerate all files found in the given list of paths. Recurse into directories as necessary.

[path list]: Newline-separated list of files and directories.

[nonexisting]: 0 = Error out with B_E on non-existing files (default). 1 =

Silently drop non-existing files. 2 = Include non-existing files in the output.

returns: List of files found in all of the given paths.

@B_E

hash

Hash functions.

Copyright (C) 2020 David Hobach LGPLv3

0.5

Dependencies

no dependencies

Imports

no imports

b_hash_file [file] [algorithm]

hash/b_hash_file

Compute the hash of the given file.

[file]: Full path to the file.

[algorithm]: Algorithm to use. Currently supported: md5|sha1|sha224|sha256|sha384|sha512|crc|blake2 (default: md5).

returns: The hash of the given file.

@B_E

b_hash_str [string] [algorithm]

hash/b_hash_str

Compute the hash of the given string.

[string]: String which to compute the hash for.

[algorithm]: Algorithm to use. Currently supported: md5|sha1|sha224|sha256|sha384|sha512|crc|blake2 (default: md5).

returns: The hash of the given string.

@B_E

http

Collection of http related functions.

Copyright (C) 2018 David Hobach LGPLv3

0.2

Dependencies

curl

Imports

types

Global Variables

B_HTTP_CHECKURLS

http/B_HTTP_CHECKURLS

Each call to `b_http_getOnlineStatus` causes one of the URLs in this array to be visited.

It is recommended to pick a relatively large number of URLs with SSL support to remain relatively anonymous, if `b_http_getOnlineStatus` is called multiple times.

Defaults to the European/US Alexa Top 10 (which hopefully blends in to the masses).

Functions

`b_http_rawUrlEncode [string]`

http/b_http_rawUrlEncode

Encode the given string according to RFC 3986.

[string]: to encode

returns: Returns a string in which all non-alphanumeric characters except `-_.~` have been replaced with a percent (%) sign followed by two hex digits. This is the encoding described in RFC 3986 for protecting literal characters from being interpreted as special URL delimiters, and for protecting URLs from being mangled by transmission media with character conversions (like some email systems). A non-zero exit code is set on errors.

@B_E

b_http_rawUrlDecode [string]

http/b_http_rawUrlDecode

Decode the given string encoded with b_str_rawUrlEncode or an equivalent function.

[string]: to decode

returns: The literal string with all hex characters replaced; a non-zero exit code is set on errors.

b_http_getOnlineStatus [timeout]

http/b_http_getOnlineStatus

Find out whether we are online or not by attempting an http connection.

One of B_HTTP_CHECKURLS is possibly visited during the process.

Use b_net_getDNSStatus for DNS-only checks.

[timeout]: Timeout in seconds for hanging checks (default: 5).

returns: 0, if we're online, 1 if only DNS works, 2 if neither DNS nor http(s) worked, 3 if the check timed out; B_E will be called if the status cannot be determined.

@B_E

b_http_testProxy [proxy string] [intransparent only] [timeout]

http/b_http_testProxy

Test whether the given proxy is working as advertised.

[proxy string]: Connection string to test, e.g. **https://1.1.1.1:234**.

[intransparent only]: If set to 0 (default), only accept intransparent proxies (i.e. those not revealing your IP).

[timeout]: Time in seconds after which to consider the proxy non-responsive (default: 5).

returns: Nothing, but sets a zero exit code, if and only if the proxy appears to work.

ini

Stateful ini reader for bash.

Currently only a single file per instance of this library/thread is kept in memory, but you can read multiple files one after another or in multiple threads.

Implementation Specifics:

- names/keys & values are case sensitive
- comment lines may start with ; or #
- whitespace lines are ignored
- duplicate names may result in undefined behaviour (usually the second will override the first)

- all characters following the = are considered part of the value (incl. whitespace); whitespace before and after the value may be trimmed by the getters though (check their description)
- values are not interpreted (e.g. quotes, escape characters, ...)
- whitespace around keys and around section qualifiers is ignored

Copyright (C) 2018 David Hobach LGPLv3
0.5

Dependencies

no dependencies

Imports

no imports

Functions

b__ini_read [ini file]

ini/b__ini_read

read the given ini file and keep it in thread-local memory so that subsequent calls to the b__ini_get functions will return the values from the ini file; subsequent calls to this function will update the internal state to represent the file last read in this thread

[ini file]: path to the ini file to read

returns: an error message on errors and sets a non-zero exit code on errors

@StateChanging

@B__E

b__ini_get [name] [section]

ini/b__ini_get

get the value for the ini entry with the given name as String in raw format

[name]: name/key of the ini entry to retrieve

[section]: section where to look for the entry with the given name (default: without section)

returns: value of the ini entry matching exactly the given section and name incl. any whitespace; a non-zero exit code is set if such an entry wasn't found

b__ini_getString [name] [section]

ini/b__ini_getString

get the value for the ini entry with the given name as String and remove all whitespace around the returned String

[name]: name/key of the ini entry to retrieve

[section]: section where to look for the entry with the given name (default:

without section)

returns: value of the ini entry matching exactly the given section and name excl. any whitespace around; a non-zero exit code is set if such an entry wasn't found

b_ini_getInt [name] [section]

ini/b_ini_getInt

get the value for the ini entry with the given name as integer

[name]: see b_ini_get

[section]: see b_ini_get

returns: see b_ini_get; additionally it is checked whether the return value is an integer (if not, a non-zero exit code of 2 is set and the return value is undefined)

b_ini_getBool [name] [section]

ini/b_ini_getBool

get the value for the ini entry with the given name as boolean

[name]: see b_ini_get

[section]: see b_ini_get

returns: see b_ini_get; 0 is returned via echo for true, 1 for false; the exit code indicates a potential error during parsing (2) or a missing entry (1) and *not* true/false

b_ini_assertNames [section 1] [name 1] ... [name n] – [section 2]
[name 1] ... [name m]

ini/b_ini_assertNames

Assert that the given sections contains *at most* the given names and no additional ones.

This is a useful function to detect user mistakes and should be called right after b_ini_read.

Multiple sections can be separated with “–”.

[section]: see b_ini_get

[name]: see b_ini_get

returns: Sets a zero exit code if and only if the last read ini file doesn't contain any additional names. Otherwise B_E is triggered. B_ERR will list the invalid names.

@B_E

keys

Simple *system-wide* cryptographic key store protected by a master password.

A key store is only opened when necessary. It'll remain open from then on until an application calls b_keys_close.

The default key store can be found at `/etc/blib/keys`. Applications may deploy their own exclusive key store, if necessary (not recommended).

Usually requires root access rights.

Features:

- thread safe: no write operation can happen to a closed key store (`b_keys_close` and write operations are single threaded only)
- retrieved keys are read-only unless managed via this interface

Copyright (C) 2020 David Hobach LGPLv3
0.6

Dependencies

cat
cp
find
findmnt
mkdir
readlink
rm

Imports

dmccrypt
fs
multithreading/mutex

Functions

`b_keys_getDefaultStore`

keys/b_keys_getDefaultStore

Get the system-wide default key store directory.

returns: The default store directory. Always sets a zero exit code.

`b_keys_init [app id] [auto create] [ui mode] [password prompt] [wait time] [store dir]`

keys/b_keys_init

Initialize this module. This function *must* be called at least once before any other function except `b_keys_create` can be used.

Will open the key store, if it's closed.

[app id]: Unique application ID (recommended: `$B_SCRIPT_NAME`).

[auto create]: Automatically create a new key store, if none exists (default: 0 = do it). Users may create their own key store using `cryptsetup` otherwise.

[ui mode]: How to request the master password from the user: `auto|gui|tty` (default: `auto`).

[password prompt]: Prompt string to ask the user for his password (optional).
 [wait time]: Maximum time in ms to wait for another process writing to a key (default: 300) (-1 = indefinitely).
 [store dir]: Full path to a directory where to manage the keys (default: `b_keys_getDefaultStore`). Only applications requiring an exclusive key store should use a non-default value here.
returns: Sets a zero exit code on success and errors out with `B_E` otherwise.
@StateChanging
@B_E

b_keys_add [key id] [key path] [keep flag]

keys/b_keys_add
 Add the given key to the key store.
 [key id]: Unique identifier for the key to add.
 [key path]: Full path to the key to add.
 [keep flag]: Keep the [key path] file after a successful copy operation (default: 0 / keep). If set to 1, it is safely removed from disk.
returns: Sets a zero exit code on success and errors out with `B_E` otherwise.
@B_E

b_keys_get [key id]

keys/b_keys_get
 Retrieve the given key from the store. If you need the content as String, please use `b_keys_getContent`.
 [key id]: Unique identifier for the key to retrieve.
returns: Path to the key. It may not exist, if the key store was closed in the meantime or the provided ID does not exist.
@B_E

b_keys_getAll [global]

keys/b_keys_getAll
 Retrieve all keys from the store.
 [global]: If set to 0, also retrieve keys for other application IDs (default: 1).
returns: All key paths for reading as a newline-separated list. If the key store is closed or modified during the runtime, the list may be empty or incomplete.
@B_E

b_keys_getContent [key id]

keys/b_keys_getContent
 Retrieve the given key content from the store. For binary keys please use `b_keys_get`.
 [key id]: Unique identifier for the key to retrieve.
returns: The key content as String. If it doesn't exist or other errors occurred,

B_E is called.

@B_E

b_keys_delete [key id] [backup]

keys/b_keys_delete

Delete the given key from the key store.

[key id]: Unique identifier for the key to delete.

[backup]: Whether (0) or not (1) to create a backup (default: 0).

returns: Sets a zero exit code on success and errors out with B_E otherwise.

@B_E

b_keys_close [store dir]

keys/b_keys_close

Close the key store, if necessary. Works without initialisation (which might open the key store), if necessary.

WARNING: 99% of all applications should *never* call this function as it may block all subsequent read operations by any other applications using this system-wide key store. The only applications that might call this function are key store management applications and only upon explicit user request. The key store is otherwise automatically closed upon system shutdown by the running OS.

[store dir]: The main directory of the store to close (default: the key store directory used in b_keys_init or - if not initialised - b_keys_getDefaultStore).

returns: Sets a zero exit code on success (closed key store) and errors out with B_E otherwise.

@B_E

meta

Functions providing information about blib modules or other bash scripts.

Copyright (C) 2020 David Hobach LGPLv3

0.2

Dependencies

no dependencies

Imports

str

Functions

b_meta_getClearImports [path]

meta/b_meta_getClearImports

Get the list of modules that the given script imports via b_import.

Important: All `b_import` calls are assumed to be declared on a dedicated line and to start at the first character of that line (“clear” import). This is considered a feature to evade this method (for doc purposes mostly).

[path]: Full path to the script to retrieve the module imports for.

returns: Newline-separated list of imports of the given script.

@B_E

b_meta_getClearDeps [path]

meta/b_meta_getClearDeps

Get the list of dependencies that the given script declares via `b_deps`.

Important: All `b_deps` calls are assumed to be declared on a dedicated line and to start at the first character of that line (“clear” dependency). This is considered a feature to evade this method (for doc purposes mostly).

[path]: Full path to the script to retrieve the dependencies for.

returns: Newline-separated list of dependencies of the given script.

@B_E

multithreading/ipcm

An inter-process map implementation.

The map is available to all running processes during a single boot session.

Both reading and writing to the map can be done from any number of processes.

Overall Features	
# readers	multiple
# writers	multiple
read consistency	always
write consistency	always
blocking	on writes

Copyright (C) 2019 David Hobach LGPLv3

0.2

Dependencies

no dependencies

Imports

multithreading/ipcv

multithreading/mtx

Functions

b_ipcm_setNamespace [namespace]

multithreading/ipcm/b_ipcm_setNamespace

Set the common process namespace to use. All processes inside the same namespace share a common state.

It is recommended to call this function a single time before any other functions of this module. Otherwise the default namespace, which may include unrelated processes, is used.

[namespace]: Name of the namespace to set.

returns: Errors out, if the name is unacceptable.

@StateChanging

@B_E

b_ipcm_getNamespace

multithreading/ipcm/b_ipcm_getNamespace

Retrieve the currently used namespace.

returns: The currently used namespace.

b_ipcm_change [key] [change function] [maximum time]

multithreading/ipcm/b_ipcm_change

Change the given key/value combination inside the map in a thread-safe way.

May wait for changes done by other processes.

[key]: A global unique identifier for the given value.

[change function]: Name of the function to execute the change of the value. It will be called with the current key as first parameter and the current value as second. It is expected to print the new value to set for this key. A non-zero exit code will cause b_ipcm_change to abort the change.

[maximum time]: maximum time in ms to wait for other processes to complete their operation (default: -1 = indefinitely)

returns: A zero exit code and prints the new value, if the change succeeded. An exit code of B_RC+1 indicates that the change function returned a non-zero exit code. B_E is called otherwise.

@B_E

b_ipcm_get [key] [fallback]

multithreading/ipcm/b_ipcm_get

Retrieve the data found at the given key.

[key]: A global unique identifier for the data to retrieve.

[fallback]: Data to return if nothing was found for the given key (default: empty).

returns: Sets a zero exit code and returns the data found on success. If no data was found, the fallback data is returned and a zero exit code is set. B_E is called on errors.

b_ipcm_unsetNamespace [namespace] [maximum time]

multithreading/ipcm/b_ipcm_unsetNamespace

Unsets the given namespace and all keys stored within it.

It is recommended to call this function when all processes finished their work.

[namespace]: The namespace to unset (default: the current namespace).

[maximum time]: maximum time in ms to wait for other processes to complete their operation (default: 0 = indefinitely)

returns: Sets a zero exit code only upon successful removal. Otherwise B_E is triggered.

@B_E

multithreading/ipcv

Provides means for inter-process communication (ipc) via global bash variables (v).

This implementation uses shared memory, i.e. it should be reasonably fast.

Only a single process or thread is assumed to be writing (i.e. use b_ipcv_save) a variable at a time and multiple processes may read it (using e.g. b_ipcv_load).

If you need to write a single variable from multiple processes, please consider using the multithreading/mtx module or similar locking means in combination with this module (or just multithreading/ipcm).

Overall Features	
# readers	multiple
# writers	single
read consistency	always
write consistency	only for one writer
blocking	never

Copyright (C) 2018 David Hobach LGPLv3

0.3

Dependencies

findmnt

mkdir

mktemp

mv

rm

Imports

no imports

Functions

b_ipcv_save [namespace] [var name 1] .. [var name n]

multithreading/ipcv/b_ipcv_save

Save the current values of the given variables so that they are made available for other processes under the given namespace.

Please note that each variable is saved atomically, but individually. I.e. if you need multiple values to be updated at the same time, please use a single variable (e.g. a map).

[namespace]: Name for a common group under which the given variables should be saved. The combination of [namespace] and [variable name] must be a unique identifier across all processes running on the system.

[var name i]: The name of the global variable to make accessible for other processes. An arbitrary number of variable names can be specified.

returns: Nothing, but a non-zero exit code indicates failed variable save attempts. A failed save attempt also triggers B_E.

@B_E

b_ipcv_load [namespace] [var name 1] .. [var name n]

multithreading/ipcv/b_ipcv_load

Load the given variables from the given namespace into the current process context.

[namespace]: Name of the group under which the variable was saved with b_ipcv_save. Must exist.

[var name i]: Name of the variable to load. Multiple names can be specified.

returns: A non-zero exit code indicates the number of variables that could not be loaded unless some unexpected error occurred and B_E is triggered. Please note that a failed load attempt/unavailable variable does generally not trigger B_E.

@B_E

b_ipcv_loadNamespace [namespace] [check existence]

multithreading/ipcv/b_ipcv_loadNamespace

Load all variables that can be loaded for the given namespace into the current process context.

[namespace]: Name of the group for which to load all available variables.

[check existence]: Whether or not to make sure that the namespace to load exists (default: 0/check). Otherwise non-existing namespaces will not cause an error.

returns: Sets a zero exit code on success. Failing to load any single available variable will always trigger B_E.

@B_E

b_ipcv_unset [namespace] [var name 1] .. [var name n]

multithreading/ipcv/b__ipcv_unset

Unset/Remove the given variables from the global namespace.

Please note that the variables will remain set in your current process context, if they were set before. Use the standard bash **unset** for that.

[namespace]: Group where the given variables belong to.

[var name i]: Name of the variable to remove. Multiple may be specified.

returns: The number of variables which could not be unset. B_E is not triggered for these.

@B_E

b__ipcv_unsetNamespace [namespace]

multithreading/ipcv/b__ipcv_unsetNamespace

Remove the given global namespace and all variables it contains.

Please note that the variables will remain set in your current process context, if they were set before. Use the standard bash **unset** for that.

[namespace]: To remove.

returns: Sets a zero exit code only upon successful removal. Otherwise B_E is triggered.

@B_E

multithreading/mtx

Collection of mutex related functions.

Mutex: Only a single process may have it at any point in time.

Semaphore: A specific maximum number of processes may have it at any point in time.

See the keys module source code for an example on how to use this mutex implementation efficiently.

Copyright (C) 2018 David Hobach LGPLv3

0.3

Dependencies

cat

mkdir

mktemp

rm

rmdir

sleep

touch

Imports

proc

Functions

b__mtx_setSleepTime [ms]

multithreading/mtx/b__mtx_setSleepTime

Sets the time to sleep for this module whenever active polling is done (default: 500).

[ms]: time in miliseconds between active polling requests for e.g. mutexes done by this module; must be an integer

returns: Nothing, always sets a zero exit code.

b__mtx_getSleepTime

multithreading/mtx/b__mtx_getSleepTime

Gets the time to sleep for this module whenever active polling is done.

returns: The currently set time to sleep in ms.

b__mtx_create [base dir]

multithreading/mtx/b__mtx_create

Allocate a new mutex without claiming it (use b__mtx_try for that).

[base dir]: Path to an *existing* directory where to store the mutex (default: not specified). By default this module will pick a temporary location. If you need a mutex that persists across reboots, please set a directory that persists across reboots here. The path should point to a local, non-network file system destination. The module must be able to create remove files or directories there at will.

returns: A string identifying the mutex (mutex ID). Sets a non-zero exit code on errors.

@B_E

b__mtx_release [mutex] [block ID]

multithreading/mtx/b__mtx_release

Release the given mutex so that it can be used by other block IDs/threads.

[mutex]: A mutex obtained via b__mtx_create.

[block ID]: The block ID for which to release the mutex (default: \$\$).

returns: Sets a non-zero exit code if the mutex could not be removed as another process is blocking it and a zero exit code on successful removal.

b__mtx_forceRelease [mutex]

multithreading/mtx/b__mtx_forceRelease

Release the given mutex so that it can be used by other blockIDs/threads.

Warning: This function can remove mutexes from other threads and should generally *only* be used for the removal of mutexes which are known to be stale by the calling application.

[mutex]: A mutex obtained via `b_mtx_create`.

returns: Nothing and sets a zero exit code.

b_mtx_pass [mutex] [block ID]

multithreading/mtx/b_mtx_pass

Pass a blocked mutex to another block ID (i.e. change the block ID of the given mutex).

You should only do this if you currently own the mutex and the new process is ready to take over.

[mutex]: A mutex obtained via `b_mtx_create`.

[block ID]: The block ID to set for the given mutex.

returns: Sets a zero exit code on success and a non-zero exit code otherwise.

@B_E

b_mtx_try [mutex] [block ID] [claim stale] [claim own]

multithreading/mtx/b_mtx_try

Attempt to obtain the given mutex. Return immediately even if it cannot be obtained.

[mutex]: A mutex obtained via `b_mtx_create`. You may also use a static and otherwise unused directory path as mutex and share it across all relevant processes.

[block ID]: The ID to use by which to block (default: running (sub)shell process id \$). This should be the process ID of the process attempting to obtain the mutex or you should know what you're doing. If you're in a subshell that should deploy a mutex against other subshells, store their \$BASHPID and call the function with that.

[claim stale]: If set to 0, claim the mutex even if it is still blocked by some other process, but that process isn't running anymore. If set to 1 (default), the function returns without obtaining the mutex. In general this should only be used in situations where a mutex has a high probability of being stale (e.g. application start).

[claim own]: If set to 0 (default), claim the mutex if it appears to be blocked by the provided block ID. If set to 1, consider it blocked even then.

returns: The function incl. parameters to execute to remove the mutex if it was obtained and an error message stating the reason otherwise. The provided function *should* be called as part of an exit trap of the calling script or via eval. Sets an exit code of 0, if the mutex was obtained. An exit code of 1 is set, if the mutex was blocked and another non-zero exit code if some other error occurred (the mutex might be blocked even then).

Example code:

```
local mutex=""
local mutexRet=""
mutex="$(b_mtx_create)" || { B_ERR="Failed to create a mutex." ; B_E }
```



```

mutexRet="$(b_mtx_try "$mutex")" \
|| { B_ERR="Failed to obtain the mutex $mutex. Reason: $mutexRet" ; B_E }
#assuming the mutex is only meant to be removed after full
#execution of the script:
trap "$mutexRet" EXIT
#direct removal:
#b_mtx_release "$mutex"

```

b_mtx_waitFor [mutex] [block ID] [claim stale] [maximum time]

multithreading/mtx/b_mtx_waitFor

Wait for the given mutex to become available. This will block script execution.

[mutex]: see b_mtx_try

[block ID]: see b_mtx_try

[claim stale]: see b_mtx_try

[maximum time]: maximum time in ms to wait for the mutex to become available
(default: -1 = indefinitely)

returns: see b_mtx_try

multithreading/multiw

Allow multiple processes to write to a *virtual* file at the same time without causing write inconsistencies (written data from each process mangled with each other).

This is achieved by keeping one file per process and relies on the assumption that both replacing and reading a symlink on your Linux distribution is atomic.

In order for this to work, all write operations must go through this module.

Currently reading only returns the data written by the process which wrote last. If you need some sort of appending, it makes more sense to deploy a mutex using e.g. the multithreading/mtx module.

Overall Features	
# readers	multiple
# writers	multiple
read consistency	partial, last writer wins
write consistency	always
blocking	never

Copyright (C) 2019 David Hobach LGPLv3

0.2

Dependencies

ln
mv
rm
shuf
stat

Imports

no imports

Functions

b__multiw__setMaxHangTime [seconds]

multithreading/multiw/b__multiw__setMaxHangTime

Set the maximum time that a process is expected to hang between two instructions. This is relevant for various internal guarantees.

[seconds]: Time in seconds that a process hangs at most.

returns: Nothing.

@StateChanging

b__multiw__getMaxHangTime

multithreading/multiw/b__multiw__getMaxHangTime

Get the maximum time that a process is expected to hang between two instructions. This is relevant for various internal guarantees.

returns: Time in seconds.

b__multiw__write [file path]

multithreading/multiw/b__multiw__write

Write all data lying in stdin to the given *virtual* file in a thread-safe way.

[file path]: Full path to the virtual file to write to. Must not be a regular file (but may not exist).

returns: A zero exit code, if the write operation was successful and a non-zero exit code otherwise.

@B_E

b__multiw__remove [file path]

multithreading/multiw/b__multiw__remove

Remove the given virtual file *and all of its revisions*.

This function should only be called when all processes finished reading and writing. It is recommended to use it over the standard Linux `rm` as the latter will leave remnants behind.

returns: A zero exit code on success.

@B_E

net

Collection of network related functions.

Copyright (C) 2020 David Hobach LGPLv3
0.1

Dependencies

dig

Imports

no imports

Global Variables

B_NET_CHECKHOSTS

net/B_NET_CHECKHOSTS

Each call to `b_net_getDNSStatus` causes one of the host records in this array to be requested.

It is recommended to pick a relatively large number of URLs with SSL support to remain relatively anonymous, if `b_net_getDNSStatus` is called multiple times. Defaults to the European/US Alexa Top 10 (which hopefully blends in to the masses).

B_NET_DNSSERVERS

net/B_NET_DNSSERVERS

Each call to `b_net_getDNSStatus` with a **random** [server] causes one of the IPs in this array to be used as DNS server.

Defaults to common large-scale free DNS providers (that probably save your requests).

Sources:

- <https://www.techradar.com/news/best-dns-server>
- <https://dnsmap.io/articles/most-popular-dns-servers>

Functions

b_net_getDNSStatus [timeout] [server]

net/b_net_getDNSStatus

Find out whether DNS appears to work or not by testing it.

One of `B_NET_CHECKHOSTS` is possibly requested during the process.

Use `b_http_getOnlineStatus` for http checks.

[timeout]: Timeout in seconds for hanging checks (default: 5).
[server]: DNS server to use (default: the OS default). Passing **random** will pick a random DNS server.
returns: 0 if DNS works as expected, 1 if DNS works but returns NXDOMAIN and 2 on timeout; B_E will be called if the status cannot be determined.
@B_E

notify

Collection of notification related functions.

Copyright (C) 2021 David Hobach LGPLv3
0.5

Dependencies

getent
notify-send
su

Imports

fs

Functions

b_notify_waitForUserDBus [user] [maximum time]

notify/b_notify_waitForUserDBus

Wait for *some* dbus user session to come up at `/run/user/[uid]/bus`. As dbus is required for notifications to work, waiting may be required for early notifications.

[maximum time]: Maximum time in s to wait for the file to appear (default: forever).

[user]: Name of the user whose instance to wait for (default: first user with a readable home directory).

returns: Sets a zero exit code and returns the identified dbus instance, if a dbus instance came up and a non-zero exit code on timeout. B_E is used for exceptional errors.

@B_E

b_notify_send [arg 1] ... [arg n]

notify/b_notify_send

Send out a notification to the user via **notify-send**.

If run as root, may send the notification to *all* active users.

Calling **notify-send** directly can be problematic depending on the user, dbus session, environment variables, ... This function aims to circumvent these

potential issues.

[arg i]: All arguments are directly passed to **notify-send**.

returns: Nothing. Errors during notification sending will trigger B_E.

@B_E

b_notify_sendNoError [arg 1] ... [arg n]

notify/b_notify_sendNoError

Convenience variant of **b_notify_send** that does not exit and only prints errors with **b_defaultErrorHandler**. This function should e.g. be used in custom error handlers to avoid recursion.

returns: Sets a non-zero exit code on errors.

os/osid

Functions for operating system identification.

Copyright (C) 2018 David Hobach LGPLv3

0.2

Dependencies

no dependencies

Imports

no imports

Functions

b_osid_init [force]

os/osid/b_osid_init

[force]: if set to 0, force an init even if it would otherwise not be necessary (default: 1 - only initialize if it didn't happen before)

Initialize the osid module. It should normally *not* be necessary to call this function directly, but it will be called by the osid module internally as needed.

returns: May error out and set a non-zero exit code on failures.

b_osid_isDebian

os/osid/b_osid_isDebian

Check whether the OS running this function is a Debian Linux.

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isDebianLike

os/osid/b__osid__isDebianLike

Check whether the OS running this function is a Debian Linux or one of its derivatives (e.g. ubuntu).

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isOpenSuse

os/osid/b__osid__isOpenSuse

Check whether the OS running this function is a OpenSUSE.

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isFedora

os/osid/b__osid__isFedora

Check whether the OS running this function is a Fedora Linux.

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isCentOS

os/osid/b__osid__isCentOS

Check whether the OS running this function is a CentOS.

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isRedHat

os/osid/b__osid__isRedHat

Check whether the OS running this function is a RedHat Linux.

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isUbuntu

os/osid/b__osid__isUbuntu

Check whether the OS running this function is an Ubuntu Linux.

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isFedoraLike

os/osid/b__osid__isFedoraLike

Check whether the OS running this function is a Fedora Linux or one of its derivatives (e.g. CentOS, Red Hat, Qubes OS).

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isQubesDom0

os/osid/b__osid__isQubesDom0

Check whether the OS running this function is a Qubes OS in dom0.

returns: Sets a zero exit code if the check returns true. Does not print any output.

b__osid__isQubesVM

os/osid/b__osid__isQubesVM

Check whether the OS running this function is a Qubes OS in a VM.

returns: Sets a zero exit code if the check returns true. Does not print any output.

os/qubes4/dom0

Collection of functions supporting scripting in Qubes OS 4.x dom0.

Important: Whenever you parse output from VMs to dom0, you **must** be extra careful and assume it totally untrusted as parsing bugs are a plausible attack vector for compromised VMs. Passing data to potentially compromised VMs of course also exposes that data's confidentiality.

Copyright (C) 2021 David Hobach LGPLv3
0.5

Dependencies

basename
cat
dd
dirname
kill
losetup
mktemp
ps
python3
qubesdb-rm
qubesdb-write
qubes-prefs
qvm-block
qvm-check
qvm-firewall
qvm-ls
qvm-pool

qvm-run
qvm-shutdown
readlink
setsid
sleep
sort
tar
tee
timeout
xxd

Imports

fs
proc
types
ui

Functions

b_dom0_setVMDeps [list] [keep defaults]

os/qubes4/dom0/b_dom0_setVMDeps

Set the dependencies for all VMs used with this module to the given list of binaries.

These dependencies are checked before *any* execution of commands in *any* VM; in particular whenever `b_dom0_qvmRun`, `b_dom0_execIn`, `b_dom0_execStrIn` or `b_dom0_execFuncIn` are called.

[list]: Newline-separated list of binaries or commands that VMs must be able to execute.

[keep defaults]: If set to 0 (default), merge the default list of dependencies with the supplied one. Otherwise only keep the supplied list of dependencies.

returns: Nothing.

@StateChanging

b_dom0_getVMDeps

os/qubes4/dom0/b_dom0_getVMDeps

Getter for `b_dom0_setVMDeps`.

returns: The current list of dependencies for VMs.

b_dom0_qvmRun [parameter 1] ... [parameter n]

os/qubes4/dom0/b_dom0_qvmRun

A wrapper for `qvm-run` which sets reasonable defaults for shell scripting and applies various fixes.

Most calls to `qvm-run` should be made via this function rather than interacting with `qvm-run` directly as the Qubes OS `qvm-run` was designed with interactive shell usage in mind whereas this wrapper is intended for bash developers.

Particular features:

- a certain set of reasonable default parameters is used: `-p -q -n -u root`
- `-n` was set as auto-starting VMs during Bash scripting can heavily influence the user experience (imagine the VM being shut down manually by the user whilst a bash script is running -> constant restarts)
- `stdin` is redirected to `/dev/null` by default to avoid potential security implications (accidental reads from `dom0` `stdin` passed to a VM); this can be overridden using `-stdin`
- `stdout` has the VM output and the exit code is the one of the VM
- distinguished exit conditions (executed command failed vs. `qvm-run` failed)
- workarounds for known Qubes bugs wrt `qvm-run` may be implemented here (e.g. qubes issues #3083, #4476, #4633 in the past)
- VM dependencies set with `b_dom0_setVMDeps` are checked

Please note that calling this function will make your script wait for the execution of the commands in the client VM.

Wherever possible, this function should be combined with `b_silence` as the VM output shouldn't be trusted. Otherwise please keep in mind that **both** `stdout` and `stderr` may have untrusted output which may even contain binary data. In order to validate against binary data you can e.g. use `b_types_parseString`.

[parameters]: Any parameters supported by `qvm-run`. If you pass `-a`, the default `-n` will be overridden. If you pass `-u`, the default root user is overridden. If you pass `-v`, `-q` will be overridden. If you pass `-stdin`, even `stdin` is passed to `qvm-run`. `-p` can be overridden by using `/dev/null` redirection. Only the short parameter versions are supported.

returns: Sets the exit code of `qvm-run` and prints its output. May error out using `B_E` if `qvm-run` itself fails.

@B_E

b_dom0_getDispVMs

os/qubes4/dom0/b_dom0_getDispVMs

Get a list of all currently existing disposable VMs.

returns: The currently existing disposable VMs as newline-separated list.

@B_E

b_dom0_startDispVM [template]

os/qubes4/dom0/b_dom0_startDispVM

Start a dispVM from the given template in the background and return its name. The disposable VM will remain started until it is shut down. If you only wish to execute a single command, please use `b_dom0_qvmRun` with the `-dispVM`

parameter.

It may take a while for this function to obtain the name of the dispVM.

[template]: The template to use for the dispVM. If no template is specified, use the default Qubes template.

returns: Name of the dispVM that was started and sets a zero exit code on success. This function may error out.

@B_E

b_dom0_execIn [vm] [file] [user]

os/qubes4/dom0/b_dom0_execIn

Execute the file as bash code in the given VM and wait for it to finish.

See b_dom0_qvmRun for various notes and words of caution.

[vm]: Name of the VM where to execute the given string. The VM is assumed to be started.

[file]: Bash file to execute in the given VM.

[user]: user as which to execute the bash file (default: root)

returns: Whatever the executed Bash code prints in the VM to stderr or stdout; the status code is set to the one of the executed Bash code on success (0). Non-zero exit codes and error messages may come from both this function as well as the code executed in the given VM.

@B_E

b_dom0_execStrIn [vm] [string] [user]

os/qubes4/dom0/b_dom0_execStrIn

Execute the String as bash code in the given VM and wait for it to finish.

Convenience wrapper for b_dom0_execIn.

See b_dom0_qvmRun for various notes and words of caution.

[vm]: see b_dom0_execIn

[string]: Bash String to execute in the given VM.

[user]: see b_dom0_execIn

returns: see b_dom0_execIn; B_E is *not* called if the executed command returns an error

@B_E

b_dom0_execFuncIn [vm] [user] [function] [module dep 1] .. [module dep n] - [function dep 1] .. [function dep d] - [function param 1] .. [function param p]

os/qubes4/dom0/b_dom0_execFuncIn

Execute the Bash function in the given VM and wait for it to finish.

Convenience wrapper for b_dom0_execIn.

See `b_dom0_qvmRun` for various notes and words of caution.

[vm]: see `b_dom0_execIn`

[user]: see `b_dom0_execIn`

[function]: Name of the function as it is declared in the current scope.

[module dep i]: Names of the modules required by the function. They do not need to be imported by the function itself.

[-]: A dash as separator character between the various parameters.

[function dep j]: An arbitrary number of functions that need to be added in order to satisfy the dependencies of the function to call (e.g. if function A is meant to be called, but uses function B internally, you'll have to pass B as one of its dependencies). Dependencies that can be found in added modules must *not* be added.

[function param p]: An arbitrary number of function parameters.

returns: see `b_dom0_execIn`; `B_E` is *not* called if the executed command returns an error

@B_E

b_dom0_waitForFileIn [vm] [file] [maximum time]

os/qubes4/dom0/b_dom0_waitForFileIn

Convenience wrapper for `b_fs_waitForFile`.

[vm]: VM where to execute.

[file]: See `[b_fs_waitForFile]`.

[maximum time]: See `[b_fs_waitForFile]`.

returns: See `[b_fs_waitForFile]`.

@B_E

b_dom0_isMountedIn [vm] [device]

os/qubes4/dom0/b_dom0_isMountedIn

Check whether the device is mounted in the given VM.

[vm]: VM where to execute.

[device]: Full path to the device (incl. `/dev/`) to check.

returns: Sets a zero exit code if the device is mounted in the VM; a non-zero exit code means that it's either not mounted or some other error occurred.

@B_E

b_dom0_mountIfNecessary [vm] [device] [mount point] [enforce]
[mount options]

os/qubes4/dom0/b_dom0_mountIfNecessary

Mount the given device in the target VM if it isn't already mounted there. Actually a wrapper for `b_fs_mountIfNecessary`.

[vm]: VM where to execute.

[device]: Full path to the device (incl. `/dev/`) to mount.

[mount point]: Full path where to mount the device. If no mount point is specified, a `/tmp/` mount point is chosen. Non-existing directories are created.

Is ignored if another mount point already exists.

[enforce]: If set to 0, enforce the given mount point to be used in addition to potentially existing ones (default: 1).

[mount options]: Options to pass to `mount`, if it needs to be executed (default: none).

returns: The chosen mount point or a newline-separated list of existing mount points on success; sets a non-zero exit code on failure. As these strings are returned from the VM, extra care must be taken when parsing them.

@B_E

b_dom0_createLoopDeviceIfNecessary [vm] [file]

os/qubes4/dom0/b_dom0_createLoopDeviceIfNecessary

Create a loop device for the file in the given VM if no old one exists. Actually a wrapper for `b_fs_createLoopDeviceIfNecessary`.

This usually requires root privileges.

[vm]: VM where to execute.

[file]: File for which to create a loop device.

returns: Created loop device or previously used one (incl. `/dev/`). Sets a non-zero exit code, if no device could be created.

@B_E

b_dom0_removeUnusedLoopDevice [vm] [device|file] [type] [map]

os/qubes4/dom0/b_dom0_removeUnusedLoopDevice

Remove a loop device from the given VM, if it is *neither* used by Qubes OS *nor* by the given VM.

Important: Calling this function may cause the device to be removed as soon as it becomes unused by the VM (and the VM *only!*). So if you called this function, you should *not* use the device any further from dom0 with e.g. `qvm-block` afterwards - even if it wasn't removed.

[vm]: VM where to remove the loop device.

[device|file]: Full path (incl. `/dev/`) to the loop device or to the backing file inside the VM.

[type]: Whether the given path represents a loop device (0) or the backing file (1). If not specified, a device will be assumed, if and only if the path starts with `/dev/`.

[map]: Optional output from a previous call to `b_dom0_parseQvmBlock`. If none is specified, this function will internally call `b_dom0_parseQvmBlock`.

returns: Sets a zero exit code, if the loop device was not used and could thus be successfully removed *or* does not exist. A nonzero exit code indicates that the device is still being used. B_E is called on other errors.

@B_E

b_dom0_copy [dom0 file] [target VM] [target VM dir] [overwrite] [parent dir]

os/qubes4/dom0/b_dom0_copy

Grab a file or directory in dom0 and push it to the given file path in the target VM.

[dom0 file]: location of the dom0 file or directory to read from, assumed to exist

[target VM]: VM to write to, assumed to exist. Must be started.

[target VM dir]: full path to the parent directory in the target VM to copy the file or directory to; non-existing parent directories are created; the name is taken from the name of the file/directory in dom0

[overwrite]: Whether or not to overwrite an existing [destination file] (default: 0 = overwrite).

[parent dir]: Set this to 0, if the [target VM dir] is the target *parent* directory (default) and to 1 if it includes the target file or folder *name* as last element.

returns: Sets an exit code of 0, if everything went fine, and a non-zero exit code otherwise.

@B_E

b_dom0_crossCopy [source VM] [source file] [target VM] [target VM dir] [overwrite] [parent dir]

os/qubes4/dom0/b_dom0_crossCopy

Cross copy a file or directory from one VM to another, initiated by dom0. No user prompt is displayed.

[source VM]: Where to copy the source file from. Must be started.

[source file]: The file or directory to copy.

[target VM]: Where to copy the file to. Must be started.

[target VM dir]: full path to the parent directory in the target VM to copy the file or directory to; non-existing parent directories are created; the name is taken from the name of the file/directory in dom0

[overwrite]: Whether or not to overwrite an existing [destination file] (default: 0 = overwrite).

[parent dir]: Set this to 0, if the [target VM dir] is the target *parent* directory (default) and to 1 if it includes the target file or folder *name* as last element.

returns: Sets an exit code of 0, if everything went fine, and a non-zero exit code otherwise.

@B_E

b_dom0_ensureRunning [vm 1] ... [vm n]

os/qubes4/dom0/b_dom0_ensureRunning

Start or unpause the given VMs if needed.

[vm i]: The VM to start.

returns: Nothing. B_E will be called for VMs that could not be started (maybe they don't exist?).

@B_E

b_dom0_isRunning [vm 1] ... [vm n]

os/qubes4/dom0/b_dom0_isRunning

Check whether the given VMs are running *and* fully operational / not hanging / not paused / not booting.

In contrast **qvm-check --running [vm]** returns true for VMs which are currently booting and **qvm-ls** doesn't check whether the OS of a VM is hanging.

[vm i]: The VM to check.

returns: Nothing. B_E will be called for VMs that are not running (or don't exist?).

@B_E

b_dom0_isHalted [vm 1] ... [vm n]

os/qubes4/dom0/b_dom0_isHalted

Check whether the given VMs are halted.

This is not necessarily the inverse of **b_dom0_isRunning** as VMs may e.g. be paused.

[vm i]: The VM to check.

returns: Nothing. B_E will be called for VMs that are not halted (or don't exist?).

@B_E

b_dom0_ensureHalted [vm 1] ... [vm n]

os/qubes4/dom0/b_dom0_ensureHalted

Shut down the given VMs if needed.

[vm i]: The VM to halt.

returns: Nothing. B_E will be called for VMs that could not be halted (maybe they don't exist?).

@B_E

b_dom0_exists [vm]

os/qubes4/dom0/b_dom0_exists

Check whether the given VM exists.

[vm]: The VM to check.

returns: Sets a zero exit code, if the VM exists a non-zero exit code otherwise.

b_dom0_openCrypt [vm] [device] [mapper name] [rw flag] [mount point] [key file] [additional options] [password prompt]

os/qubes4/dom0/b_dom0_openCrypt

In the given VM, open the given crypto device with dm-crypt and mount it to the mount point.

[vm]: The VM where to open the crypto device.

[device]: Full path to the device (incl. /dev/) to open.

[mapper name]: The name to assign to the decrypted version of the crypto block device. The created decrypted device will be found at */dev/mapper/[mapper*

name].

[rw flag]: 0=open read-write, 1=open read-only (default: 0)

[mount point]: Where to mount the decrypted data to. Non-existing directories will be created. If no mount point is specified, it will not be mounted (default).

[key file]: Full vm path to the key to use for decryption. If none is specified, password-based decryption is assumed and stdin will be read to obtain the password.

[additional options]: Single string with additional cryptsetup parameters to pass on (default: none). They are passed *as-is*, i.e. you'll have to take care of proper escaping etc. yourself.

[password prompt]: Optional string to use when the user is required to provide a decryption password.

returns: nothing (except for user interaction prompts if no key file is provided), but sets a non-zero exit code on errors

@B_E

b_dom0_closeCrypt [vm] [mapper name] [mount point]

os/qubes4/dom0/b_dom0_closeCrypt

Close a crypto device opened with `b_dom0_openCrypt`.

[vm]: The VM where to close the crypto device.

[mapper name]: The name used when it was opened.

[mount point]: If the decrypted data is mounted inside the [vm], please specify the mount point here so that it can be unmounted before closing the device. Otherwise the function will attempt to close the device without unmounting (likely to fail).

returns: nothing, but sets a non-zero exit code on errors

@B_E

b_dom0_parseQvmBlock [variable name] [input]

os/qubes4/dom0/b_dom0_parseQvmBlock

Parse data from `qvm-block ls` to an associative array.

The indices of the associative array will be of the format *[counter]_[field]* (counters run from 0 (inclusive) to max (exclusive)).

The special index “max” is equal to the number of lines. It can be used for iterations over the map.

Currently supported [field] values: backend|device id|id|description|used by|read-only|frontend-dev

[variable name]: Name of the associative array to use as output.

[input]: Optional output from a previous `qvm-block ls` call; if none is specified, this function will execute the call and use its output.

returns: A string specifying an associative array in bash syntax. You can eval that string to obtain all relevant data or use the more convenient `b_dom0_getQvmBlockInfo`. On errors B_E is called.

@B_E

b_dom0_getQvmBlockInfo [map] [retrieve field] [filter field 1] [filter value 1] .. [filter field n] .. [filter value n]

os/qubes4/dom0/b_dom0_getQvmBlockInfo

Convenience function to retrieve information from the output of `b_dom0_parseQvmBlock`. Searches for the given filter values in the given fields and retrieves the first field value matching all filters.

More simple, but less flexible than `b_dom0_parseQvmBlock`.

[map]: Optional output from a previous call to `b_dom0_parseQvmBlock`. If none is specified, this function will internally call `b_dom0_parseQvmBlock`.

[retrieve field]: Name of the field to retrieve.

[filter field i]: Name of any field supported by `b_dom0_parseQvmBlock`.

[filter value i]: The value to search for in [filter field i] (equality check).

returns: The value of the requested field and sets a zero exit code on success. Sets a non-zero exit code if no matching value could be found. On errors `B_E` is called.

@B_E

b_dom0_attachDevice [dom0 device] [target VM] [rw flag] [force]

os/qubes4/dom0/b_dom0_attachDevice

Attach the given device from dom0 (!) as block device to the target VM.

[dom0 device]: Full path to the device node (incl. /dev/) in dom0.

[target VM]: VM to attach the device to. Must be started.

[rw flag]: If set to 0, attaches the dom0 device in r/w (read-write) mode. If set to 1 (default), attaches the file in r/o (read only) mode.

[force]: If set to 0, attempt to force Qubes OS to recognize the device (default: 1). This is required for e.g. VM images. Please call `b_dom0_detachDevice` with the clean flag for the detach operation, if set to 0.

returns: The full path to the device created in the target VM and sets a zero exit code on success. Otherwise a non-zero exit code is set.

@B_E

b_dom0_attachFile [dom0 file] [target VM] [rw flag] [force]

os/qubes4/dom0/b_dom0_attachFile

Attach the given file from dom0 (!) as block device to the target VM.

The function may attempt to acquire root privileges (and thus display a password prompt).

[dom0 file]: Full path to the file in dom0 to attach.

[target VM]: VM to attach the file to. Must be started.

[rw flag]: If set to 0, attaches the dom0 file in r/w (read-write) mode. If set to 1 (default), attaches the file in r/o (read only) mode.

[force]: If set to 0, attempt to force Qubes OS to recognize the device (default: 1). This is required for e.g. VM images. Please call `b_dom0_detachDevice` with the clean flag for the detach operation, if set to 0.

returns: The full path to the device created in the target VM and sets a zero

exit code on success. Otherwise a non-zero exit code is set.

@B_E

b_dom0_attachVMDisk [source VM] [target VM] [rw flag] [pool driver] [disk path]

os/qubes4/dom0/b_dom0_attachVMDisk

Attach the entire private disk image (private.img) of the source VM to the target VM.

Warning: This is contradictory to all Qubes principles and should only be done if you know exactly what you're doing. Qubes OS even has some countermeasures to prevent accidental use of this feature which are bypassed here.

[source VM]: Name of the VM whose private disk to attach to the target VM. *All data* of that VM will be shared with the target VM. Will be shut down as part of this function and must remain shut down as long as the disk is attached.

[target VM]: VM where to attach the disk as block device to. Must be started.

[rw flag]: If set to 0, attaches the disk file in r/w (read-write) mode. If set to 1 (default), attaches the file in r/o (read only) mode.

[pool driver]: Driver of the pool that the disk was created in: One of file, file-reflink or lvm_thin (default: driver of your system's default pool).

[disk path]: Full path to the disk. Default: Path to the private volume of the [source VM] in the default pool of the respective [pool type]. May override the [source VM] parameter.

returns: The full path to the device created in the target VM and sets a zero exit code on success. Otherwise a non-zero exit code is set. Please call b_dom0_detachDevice with the clean flag for the detach operation.

@B_E

b_dom0_getDefaultPoolDriver

os/qubes4/dom0/b_dom0_getDefaultPoolDriver

Retrieve the driver name of the system's default pool.

returns: Driver name of the system's default pool.

@B_E

b_dom0_crossAttachDevice [source VM] [source device] [target VM] [rw flag]

os/qubes4/dom0/b_dom0_crossAttachDevice

Attach the given block device from the source VM to the target VM.

This is merely a convenience wrapper for **qvm-block attach**.

[source VM]: VM where the source file can be found.

[source device]: Device to attach to the [target VM].

[target VM]: VM to attach the device to. Must be started.

[rw flag]: If set to 0, attaches the [source device] in r/w (read-write) mode. If set to 1 (default), attaches the file in r/o (read only) mode.

returns: The full path to the device created in the target VM and sets a zero

exit code on success. Otherwise a non-zero exit code is set.

@B_E

b_dom0_crossAttachFile [source VM] [source file] [target VM] [rw flag]

os/qubes4/dom0/b_dom0_crossAttachFile

Attach the given file from the source VM as block device to the target VM.

[source VM]: VM where the source file can be found.

[source file]: File to attach as block device.

[target VM]: VM to attach the file to. Must be started.

[rw flag]: If set to 0, attaches the [source file] in r/w (read-write) mode. If set to 1 (default), attaches the file in r/o (read only) mode.

returns: The full path to the device created in the target VM and sets a zero exit code on success. Otherwise a non-zero exit code is set.

@B_E

b_dom0_detachDevice [vm] [device] [clean flag]

os/qubes4/dom0/b_dom0_detachDevice

Attempts to detach the given device from the VM. This may fail if the VM is using the device and thus it is usually a better idea to just shut the VM down.

[vm]: VM from which to detach the device.

[device]: Full path to the device in the VM. E.g. the return values of `b_dom0_crossAttachFile`, `b_dom0_attachFile` or `b_dom0_attachVMDisk`.

[clean flag]: Also wipe the device from the Qubes DB, if it's there (default: 1).

returns: nothing, but sets a zero exit code on success

@B_E

b_dom0_enterEventLoop [callback function] [heartbeat interval]

os/qubes4/dom0/b_dom0_enterEventLoop

Enter a blocking loop to react to Qubes OS events. `b_dom0_disconnectEventLoop` can be used to disconnect from the event loop asynchronously, e.g. from exit traps.

If you want to obtain an overview of Qubes OS events, please use the `qwatch` utility manually.

[callback function]: Name of the function to call for *every* Qubes OS event. Since the number of events may be high, the function should do appropriate filtering at high performance.

The callback function should be declared as follows:

`callback_function_name` [subject] [event name] [event info] [timestamp]

[subject]: The subject name Qubes OS provides. Usually the VM for which the event was reported. 'None' appears to mean 'dom0'.

[event name]: Name of the event for which the callback function was called.

[event info]: May contain additional information about the event (e.g.

arguments).

[timestamp]: When the event was received in ms since EPOCH.

returns: Nothing. A non-zero exit code will abort further processing.

[heartbeat interval]: Interval in ms at which to send heartbeat events (default: no heartbeat events).

returns: Nothing, but uses the exit code of the last callback function execution as its own. Sets an exit code of 0, if the event loop was terminated by `b_dom0_disconnectEventLoop`. `B_E` is only called on exceptional errors.

@StateChanging

@B_E

b_dom0_disconnectEventLoop

os/qubes4/dom0/b_dom0_disconnectEventLoop

Send a termination request to a dom0 event loop started with `b_dom0_enterEventLoop`. This is useful to exit the event loop from outside of it.

returns: A non-zero exit code indicates that either no active event loop existed or the signal sending failed. Never errors out to be suitable for e.g. exit traps.

@StateChanging

b_dom0_testMultiple [vm] [operator] [list]

os/qubes4/dom0/b_dom0_testMultiple

Run multiple `test` checks inside the given VM with a single `qvm-run` call.

[vm]: Where to run the checks.

[operator]: The test operator to use for the list of parameters, e.g. `-f`.

[list]: Newline-separated list of parameters to check, e.g. a list of files for the `-f` operator.

returns: The subset of the list that matched the operator and sets a zero exit code on success.

@B_E

b_dom0_applyFirewall [vm] [rules]

os/qubes4/dom0/b_dom0_applyFirewall

Apply the given rules to the VM firewall.

Important: Since the rules are applied iteratively, there may be short timeframes during which a VM may have a more permissive firewall than in the end.

[vm]: VM for which to apply the firewall rules.

[rules]: Newline-separated list of firewall rules in *raw* format (cf. `man qvm-firewall`). The rules are added in the order as specified. (This corresponds to the `qvm-firewall` output.) Previously applied rules are left untouched.

Empty lines and lines starting with `#` (comments) are ignored.

returns: Nothing. If a rule cannot be added for whatever reason, the entire

firewall is cleared to a drop-all state for security reasons.

@B_E

b_dom0_clearFirewall [vm]

os/qubes4/dom0/b_dom0_clearFirewall

Clears the given VM firewall entirely (no rules left). This will make the VM firewall drop all connections.

Existing connections may continue to work as long as the VM and the connection remain up (this is a known Qubes OS quirk).

[vm]: VM for which to clear all rules.

returns: Nothing. Errors may result in leftover rules which should generally allow less.

@B_E

proc

Collection of process and thread related functions.

Copyright (C) 2020 David Hobach LGPLv3

0.2

Dependencies

kill

tail

timeout

Imports

no imports

Functions

b_proc_pidExists [pid]

proc/b_proc_pidExists

Check whether the given process ID exists on the system.

[pid]: process ID to check for existence (process exists)

returns: A zero exit code, if it exists and a non-zero exit code if it doesn't; this function attempts to check the existence of the given process across *all* users, but it cannot guarantee correctness if the user running this script has very low privileges.

b_proc_childExists [pid]

proc/b_proc_childExists

Check whether the given process ID exists *and* is a *direct* child of the calling

bash process.

[pid]: process ID to check

returns: A zero exit code, if the caller is a parent of the given pid.

b_proc_waitForPid [pid] [maximum time]

proc/b_proc_waitForPid

Wait for the given process to exit. If it doesn't exist, exit immediately.

[pid]: Process ID of the process to wait for. Doesn't need to be a child.

[maximum time]: Maximum time in seconds to wait for the process to exit (default: 0 = indefinitely).

returns: Nothing, always sets a zero exit code. Use `b_proc_pidExists` if you need to know whether the process finished.

b_proc_resolveSignal [signal]

proc/b_proc_resolveSignal

Resolve process (kill) signal names to their numeric identifiers.

[signal]: Signal name (string) or numeric identifier.

returns: The numeric identifier corresponding to the given signal.

@B_E

sqlite3

Stateful sqlite driver for bash.

This module provides all features of the sqlite interactive mode to non-interactive bash scripts. See `man sqlite3` for the available commands.

Keeping a single database connection open usually exhibits better performance than calling `sqlite3` in batch mode over and over again.

Side Note: If you want to read or write csv files, this module can also do the job for you with standard SQL syntax. See `man sqlite3` on how to read and write csv files with sqlite.

Copyright (C) 2020 David Hobach LGPLv3

0.5

Dependencies

mkfifo

mktemp

rm

sqlite3

Imports

no imports

Functions

b_sqlite3_open [db file] [timeout]

sqlite3/b_sqlite3_open

Open a new sqlite connection to a database. This **must** be called exactly once before any calls to `b_sqlite3_exec`.

[db file]: The database file to connect to (default: a new in-memory database). If it doesn't exist, it may be created (the behaviour is identical to `sqlite3`).

[timeout]: Maximum time in seconds to wait for a command to execute on the database via `b_sqlite3_exec`. Default: -1 = indefinitely

returns: Nothing.

@StateChanging

@B_E

b_sqlite3_getOpen

sqlite3/b_sqlite3_getOpen

Retrieve the currently open database.

returns: Prints the currently open database and an empty string for an unnamed database. Sets a zero exit code, if and only if a database is currently open.

b_sqlite3_exec [command] [timeout] [filter input]

sqlite3/b_sqlite3_exec

Executes the given command on an open sqlite database (cf. `b_sqlite3_open`).

[command]: The command to execute. All commands supported by `sqlite3` in interactive mode incl. the SQL commands are supported. The only exception is the `.output` command. Please use the bash output redirection instead of `.output`. **Warning:** Incomplete (e.g. missing `;`) or incorrect commands may cause this function to make your bash script hang forever, if no [timeout] is specified.

[timeout]: Maximum time in seconds to wait for the command to execute on the database (default: the timeout initialized via `b_sqlite3_open`).

[filter input]: Whether or not to filter command input lines from the returned output (default: 0 = true).

returns: The output provided by the database in response to the executed command incl. potential errors. The function attempts to set a non-zero exit code, if the output contains error messages. `B_E` is only called on timeouts or other database connectivity issues.

@B_E

b_sqlite3_close

sqlite3/b_sqlite3_close

Closes a currently open database connection.

It is highly recommended to execute this function once you don't need the database connection anymore (usually at the end of a script).

returns: Nothing. A non-zero exit code indicates a failed close operation.

@StateChanging

str

Collection of string related functions.

Copyright (C) 2021 David Hobach LGPLv3

0.2

Dependencies

no dependencies

Imports

no imports

Functions

b_str_stripQuotes [string]

str/b_str_stripQuotes

Remove any single or double quotes around the given string.

[string]: string which might be enclosed in single or double quotes (' or ")

returns: [string] without the enclosed single or double quotes, if there were any; if none were found the original string is returned; the exit code is always zero

b_str_trim [string]

str/b_str_trim

Remove any whitespace from around a string.

[string]: string to trim

returns: [string] beginning and ending without whitespace; the exit code is always zero

b_str_prefixLines [string] [prefix]

str/b_str_prefixLines

Prefix all lines of the given string with a given prefix.

[string]: Each line of this string will be prefixed.

[prefix]: The string to put in front of each line.

returns: All lines of the input string, each of them prefixed with the given prefix.

tcolors

Defines some tput related constants. In order to change terminal colors you can then use something such as

```
echo "$(tput setaf ${B_TCOLORS[red]})This is red, \  
$(tput setaf ${B_TCOLORS[blue]})this blue, $(tput sgr0)this normal."
```

tput can do a lot more than colors, see: man tput & man terminfo.

Copyright (C) 2018 David Hobach LGPLv3

0.1

Dependencies

tput

Imports

no imports

Global Variables

B_TCOLORS

tcOLORS/B_TCOLORS

Global map for human readable colors to tput style color identifiers.

Currently supported values: black|red|green|yellow|blue|magenta|cyan|white

traps

Collection of trap related functions.

Copyright (C) 2018 David Hobach LGPLv3

0.2

Dependencies

no dependencies

Imports

no imports

Functions

b_traps_getCodeFor [signal]

traps/b_traps_getCodeFor

Retrieve the current trap code / commands for the given signal.

returns: The current code and sets a zero exit code on success.

@B_E

b_traps_add [code] [signal] [tag] [append flag]

traps/b_traps_add

Add the given commands to the given trap signal.

[code]: Whatever should be added to the trap.

[signal]: Name of the signal to add the commands to.

[tag]: An optional *unique* marker for these commands so that they can be removed with `b_traps_remove` later on.

[append flag]: Whether to append the new commands to the end (0: default) or insert them in the beginning (1).

returns: Whatever the internal call to *trap* to set the new trap returns.

@B_E

b_traps_prepend [code] [signal] [tag]

traps/b_traps_prepend

Prepend the given commands to the ones currently existing for the given trap signal.

Convenience wrapper to `b_traps_add` with [append flag] set to 1.

[code]: see `b_traps_add`

[signal]: see `b_traps_add`

[tag]: see `b_traps_add`

returns: see `b_traps_add`

@B_E

b_traps_remove [signal] [tag]

traps/b_traps_remove

Remove the commands tagged with the given tag from the signal trap.

[signal]: Name of the signal to remove the commands from.

[tag]: The *unique* marker to identify the commands to be removed.

returns: Nothing, but sets a zero exit code on success. May error out if the tag isn't found or the internal trap call failed.

@B_E

types

Functions for data type checks and conversions.

Copyright (C) 2020 David Hobach LGPLv3

0.4

Dependencies

head

mkfifo

mktemp

rm
strings
tee
wc

Imports

no imports

Global Variables

B_TYPES_ENCODINGS

types/B_TYPES_ENCODINGS

Global map for human readable string encodings which can be used for `b_types_parseString`.

Currently supported values: 7-bit|8-bit|16-bit-bigendian|16-bit-littleendian|32-bit-bigendian|32-bit-littleendian

See the strings manpage for further explanations.

Functions

`b_types_parseString [encoding]`

types/b_types_parseString

Checks whether whatever is lying in `stdin` is a string (and not binary) and if so, prints it to `stdout`.

Important:

- bash has major issues whenever binary data is involved. For example equality checks may return undefined results. So whenever you are unsure as to whether a variable is a string or not, better pass it thorough this function.
- The input is taken from *stdin* rather than as parameter as binary parameters may also cause issues (special bytes etc.).
- Even builtins such as `echo` do not necessarily play well with binary data. So it is recommended to pipe binary data through this function before further processing.

Examples:

#check a file

```
b_types_parseString < "/path/to/potential/binary" > /dev/null && echo "It is a string file."
```

#read parts of a file as string

```
str="$(dd if="/path/to/another/file" bs=1 skip=8 | b_types_parseString)"  
[ $? -eq 0 ] && echo "Found the following string: $str"
```

[encoding]: The encoding of the string lying in stdin. Use `B_TYPES_ENCODINGS` for this parameter. Defaults to `${B_TYPES_ENCODINGS["7-bit"]}`, which makes sense in 99% of all cases as scripts should use ASCII only anyway (when no user-interaction is involved) in order to remain portable. Keep in mind that bash also needs to support the target encoding in order to support further processing.

returns: The data as String, if the input data was found to be a String. If no String was found to be lying in stdin, the output is an undefined string and a non-zero exit code is set. `B_E` is only called on exceptional errors.

@B_E

b_types_isInteger [string]

types/b_types_isInteger

Check whether the given String is an integer (positive or negative) or not.

[string]: The string to check. If it may be binary data, please make sure to pass it through `b_types_parseString` first.

returns: Nothing, but sets a zero exit code if and only if the given string represents an integer.

b_types_assertInteger [string] [error msg]

types/b_types_assertInteger

Check whether the given String is an integer (positive or negative) and if not, error out.

[string]: The string to check. If it may be binary data, please make sure to pass it through `b_types_parseString` first.

[error msg]: Error message to use, if the check fails (optional).

returns: Nothing. If it's no integer, `B_E` is called.

@B_E

b_types_looksLikeArray [string]

types/b_types_looksLikeArray

Check whether the given String "looks like" a bash array or not.

It may still contain malicious code or whatnot. So you **must not** rely on this function when processing untrusted input.

[string]: The string to check. Use `"$(declare -p var)"` on variables to obtain it.

returns: Nothing, but sets a zero exit code if the given string looks like it could be eval'ed to a bash array. References (`declare -n`) will result in a non-zero exit code.

b_types_assertLooksLikeArray [string] [error msg]

types/b_types_assertLooksLikeArray

Check whether the given String "looks like" a bash array and if not, error out.

It may still contain malicious code or whatnot. So you **must not** rely on this function when processing untrusted input.

[string]: The string to check. Use "\$(**declare -p var**)" on variables to obtain it.

[error msg]: Error message to use, if the check fails (optional).

returns: Nothing. If doesn't look like a bash array, B_E is called.

@B_E

b_types_looksLikeMap [string]

types/b_types_looksLikeMap

Check whether the given String "looks like" a bash map / associative array or not.

It may still contain malicious code or whatnot. So you **must not** rely on this function when processing untrusted input.

[string]: The string to check. Use "\$(**declare -p var**)" on variables to obtain it.

returns: Nothing, but sets a zero exit code if the given string looks like it could be **eval**'ed to a bash associative array. References (**declare -n**) will result in a non-zero exit code.

b_types_assertLooksLikeMap [string] [error msg]

types/b_types_assertLooksLikeMap

Check whether the given String "looks like" a bash map / associative array and if not, error out.

It may still contain malicious code or whatnot. So you **must not** rely on this function when processing untrusted input.

[string]: The string to check. Use "\$(**declare -p var**)" on variables to obtain it.

[error msg]: Error message to use, if the check fails (optional).

returns: Nothing. If doesn't look like a bash array, B_E is called.

@B_E

ui

Collection of user interface and user interaction related functions.

Copyright (C) 2020 David Hobach LGPLv3

0.2

Dependencies

no dependencies

Imports

no imports

Functions

b_ui_passwordPrompt [output var] [ui mode] [prompt string]

ui/b_ui_passwordPrompt

Prompt the user for a password. Should run inside the parent process (i.e. not inside a subshell).

Once you do not need the password anymore, it is recommended to wipe it from memory as such:

```
#overwrite the password in memory with zeroes, then free it  
pass="${pass//?/0}" ; pass=""
```

[output var]: The name of the variable to write the password to.

[ui mode]: How to request the password from the user: auto|gui|tty (default: auto).

[prompt string]: The string to present to the user asking for the password (default: "Password:").

returns: Nothing. Sets a non-zero exit code on errors.

@B_E

wm

Collection of functions related to window managers.

Copyright (C) 2020 David Hobach LGPLv3

0.1

Dependencies

wmctrl

Imports

no imports

Functions

b_wm_getActiveWindowProperties [variable name]

wm/b_wm_getActiveWindowProperties

Retrieve information about the currently active windows.

[variable name]: Name of the associative array to use as output.

returns: A string specifying an associative array in bash syntax. You can eval that string to obtain all relevant data. On errors B_E is called.

The information is returned as an associative array containing the following data:

[window ID]__[id] = hexadecimal window ID

[window ID]__[desktop] = desktop number

[window ID]__[pid] = PID of the window

[window ID]__[x] = x-offset
[window ID]__[y] = y-offset
[window ID]__[width] = window width
[window ID]__[height] = window height
[window ID]__[class] = window class name
[window ID]__[client] = client machine name
[window ID]__[title] = window title
Missing data may result in an empty String for the respective property.
@B_E

Reference List

B_ARGS
b_args_assertOptions
b_args_get
b_args_getAll
b_args_getAllOptions
b_args_getCount
b_args_getInt
b_args_getOption
b_args_getOptionCount
b_args_getOptionInt
b_args_getOptionParamSeparator
b_args_init
B_ARGS_OPTS
b_args_parse
b_args_setOptionParamSeparator
b_arr_contains
b_arr_join
b_arr_mapsAreEqual
b_arr_toList
b_cachingMessageHandler
B_CALLER_NAME
b_cdoc_cbPrintFirstParam
b_cdoc_cbPrintNewline

b_cdoc_generate
b_cdoc_generateBlibStyle
b_cdoc_getBlockCallback
b_cdoc_getDocumentBeginCallback
b_cdoc_getDocumentEndCallback
b_cdoc_getExtractionRegex
b_cdoc_getFileCallback
b_cdoc_getPostProcessingCallback
b_cdoc_setBlockCallback
b_cdoc_setDocumentBeginCallback
b_cdoc_setDocumentEndCallback
b_cdoc_setExtractionRegex
b_cdoc_setFileCallback
b_cdoc_setPostProcessingCallback
b_checkVersion
B_CONF_DIR
b_daemon_getPid
B_DAEMON_ID
b_daemon_init
b_daemon_restart
b_daemon_sendSignal
b_daemon_start
b_daemon_status
b_daemon_statusPid
b_daemon_stop
b_date_add
b_date_addDays
b_date_diff
b_date_getFileModAge
b_defaultErrorHandler
b_defaultMessageHandler

b_delay_execute
B_DELAY_EXECUTED
b_delay_getCommandAt
b_delay_to
b_deps
b_dmccrypt_close
b_dmccrypt_createLuks
b_dmccrypt_getMapperName
b_dmccrypt_init
b_dmccrypt_isOpen
b_dmccrypt_open
b_dom0_applyFirewall
b_dom0_attachDevice
b_dom0_attachFile
b_dom0_attachVMDisk
b_dom0_clearFirewall
b_dom0_closeCrypt
b_dom0_copy
b_dom0_createLoopDeviceIfNecessary
b_dom0_crossAttachDevice
b_dom0_crossAttachFile
b_dom0_crossCopy
b_dom0_detachDevice
b_dom0_disconnectEventLoop
b_dom0_ensureHalted
b_dom0_ensureRunning
b_dom0_enterEventLoop
b_dom0_execFuncIn
b_dom0_execIn
b_dom0_execStrIn
b_dom0_exists

b_dom0_getDefaultPoolDriver
b_dom0_getDispVMs
b_dom0_getQvmBlockInfo
b_dom0_getVMDeps
b_dom0_isHalted
b_dom0_isMountedIn
b_dom0_isRunning
b_dom0_mountIfNecessary
b_dom0_openCrypt
b_dom0_parseQvmBlock
b_dom0_qvmRun
b_dom0_removeUnusedLoopDevice
b_dom0_setVMDeps
b_dom0_startDispVM
b_dom0_testMultiple
b_dom0_waitForFileIn
B_E
b_enforceUser
B_ERR
b_error
b_execFuncAs
b_execFuncInCurrentContext
b_flog_close
b_flog_defaultHeader
b_flog_getDateFormat
b_flog_getHeaderFunction
b_flog_getLogReductionLinesLowerBound
b_flog_getLogReductionLinesUpperBound
b_flog_headerDateScriptSeverity
b_flog_headerDateSeverity
b_flog_init

b_flog_log
b_flog_messageHandler
b_flog_printSeverity
b_flog_setDateFormat
b_flog_setHeaderFunction
b_flog_setLogReductionLinesApprox
b_flog_setLogReductionLinesLowerBound
b_flog_setLogReductionLinesUpperBound
B_FLOG_SEV
b_fs_createLoopDeviceIfNecessary
b_fs_enumerate
b_fs_getLineCount
b_fs_getMountpoints
b_fs_isEmptyDir
b_fs_mountIfNecessary
b_fs_parseSize
b_fs_removeUnusedLoopDevice
b_fs_removeWithOverwrite
b_fs_waitForFile
b_generateStandalone
b_getBlibModules
b_getDefaultMessageHandlerIntermediate
b_getDefaultMessageHandlerPrefix
b_getErrorHandler
b_getMessageHandler
b_hash_file
b_hash_str
B_HTTP_CHECKURLS
b_http_getOnlineStatus
b_http_rawUrlDecode
b_http_rawUrlEncode

b_http_testProxy
b_import
b_info
b_ini_assertNames
b_ini_get
b_ini_getBool
b_ini_getInt
b_ini_getString
b_ini_read
b_initCachingMessageHandler
b_ipcm_change
b_ipcm_get
b_ipcm_getNamespace
b_ipcm_setNamespace
b_ipcm_unsetNamespace
b_ipcv_load
b_ipcv_loadNamespace
b_ipcv_save
b_ipcv_unset
b_ipcv_unsetNamespace
b_isFunction
b_isModule
b_keys_add
b_keys_close
b_keys_delete
b_keys_get
b_keys_getAll
b_keys_getContent
b_keys_getDefaultStore
b_keys_init
B_LIB_DIR

b_listContains
b_meta_getClearDeps
b_meta_getClearImports
b_mtx_create
b_mtx_forceRelease
b_mtx_getSleepTime
b_mtx_pass
b_mtx_release
b_mtx_setSleepTime
b_mtx_try
b_mtx_waitFor
b_multiw_getMaxHangTime
b_multiw_remove
b_multiw_setMaxHangTime
b_multiw_write
B_NET_CHECKHOSTS
B_NET_DNSSERVERS
b_net_getDNSStatus
b_nop
b_notify_send
b_notify_sendNoError
b_notify_waitForUserDBus
b_osid_init
b_osid_isCentOS
b_osid_isDebian
b_osid_isDebianLike
b_osid_isFedora
b_osid_isFedoraLike
b_osid_isOpenSuse
b_osid_isQubesDom0
b_osid_isQubesVM

b_osid_isRedHat
b_osid_isUbuntu
b_printStackTrace
b_proc_childExists
b_proc_pidExists
b_proc_resolveSignal
b_proc_waitForPid
B_RC
b_resetErrorHandler
B_SCRIPT
B_SCRIPT_DIR
B_SCRIPT_NAME
b_setBE
b_setDefaultMessageHandlerIntermediate
b_setDefaultMessageHandlerPrefix
b_setErrorHandler
b_setMessageHandler
b_silence
b_sqlite3_close
b_sqlite3_exec
b_sqlite3_getOpen
b_sqlite3_open
b_str_prefixLines
b_str_stripQuotes
b_str_trim
B_TCOLORS
B_TEST_MODE
b_traps_add
b_traps_getCodeFor
b_traps_prepend
b_traps_remove

b_types_assertInteger
b_types_assertLooksLikeArray
b_types_assertLooksLikeMap
B_TYPES_ENCODINGS
b_types_isInteger
b_types_looksLikeArray
b_types_looksLikeMap
b_types_parseString
b_ui_passwordPrompt
b_version
b_wm_getActiveWindowProperties